



Instruction Set Architecture (Part 1)

Support for Programming Languages

Chester Rebeiro

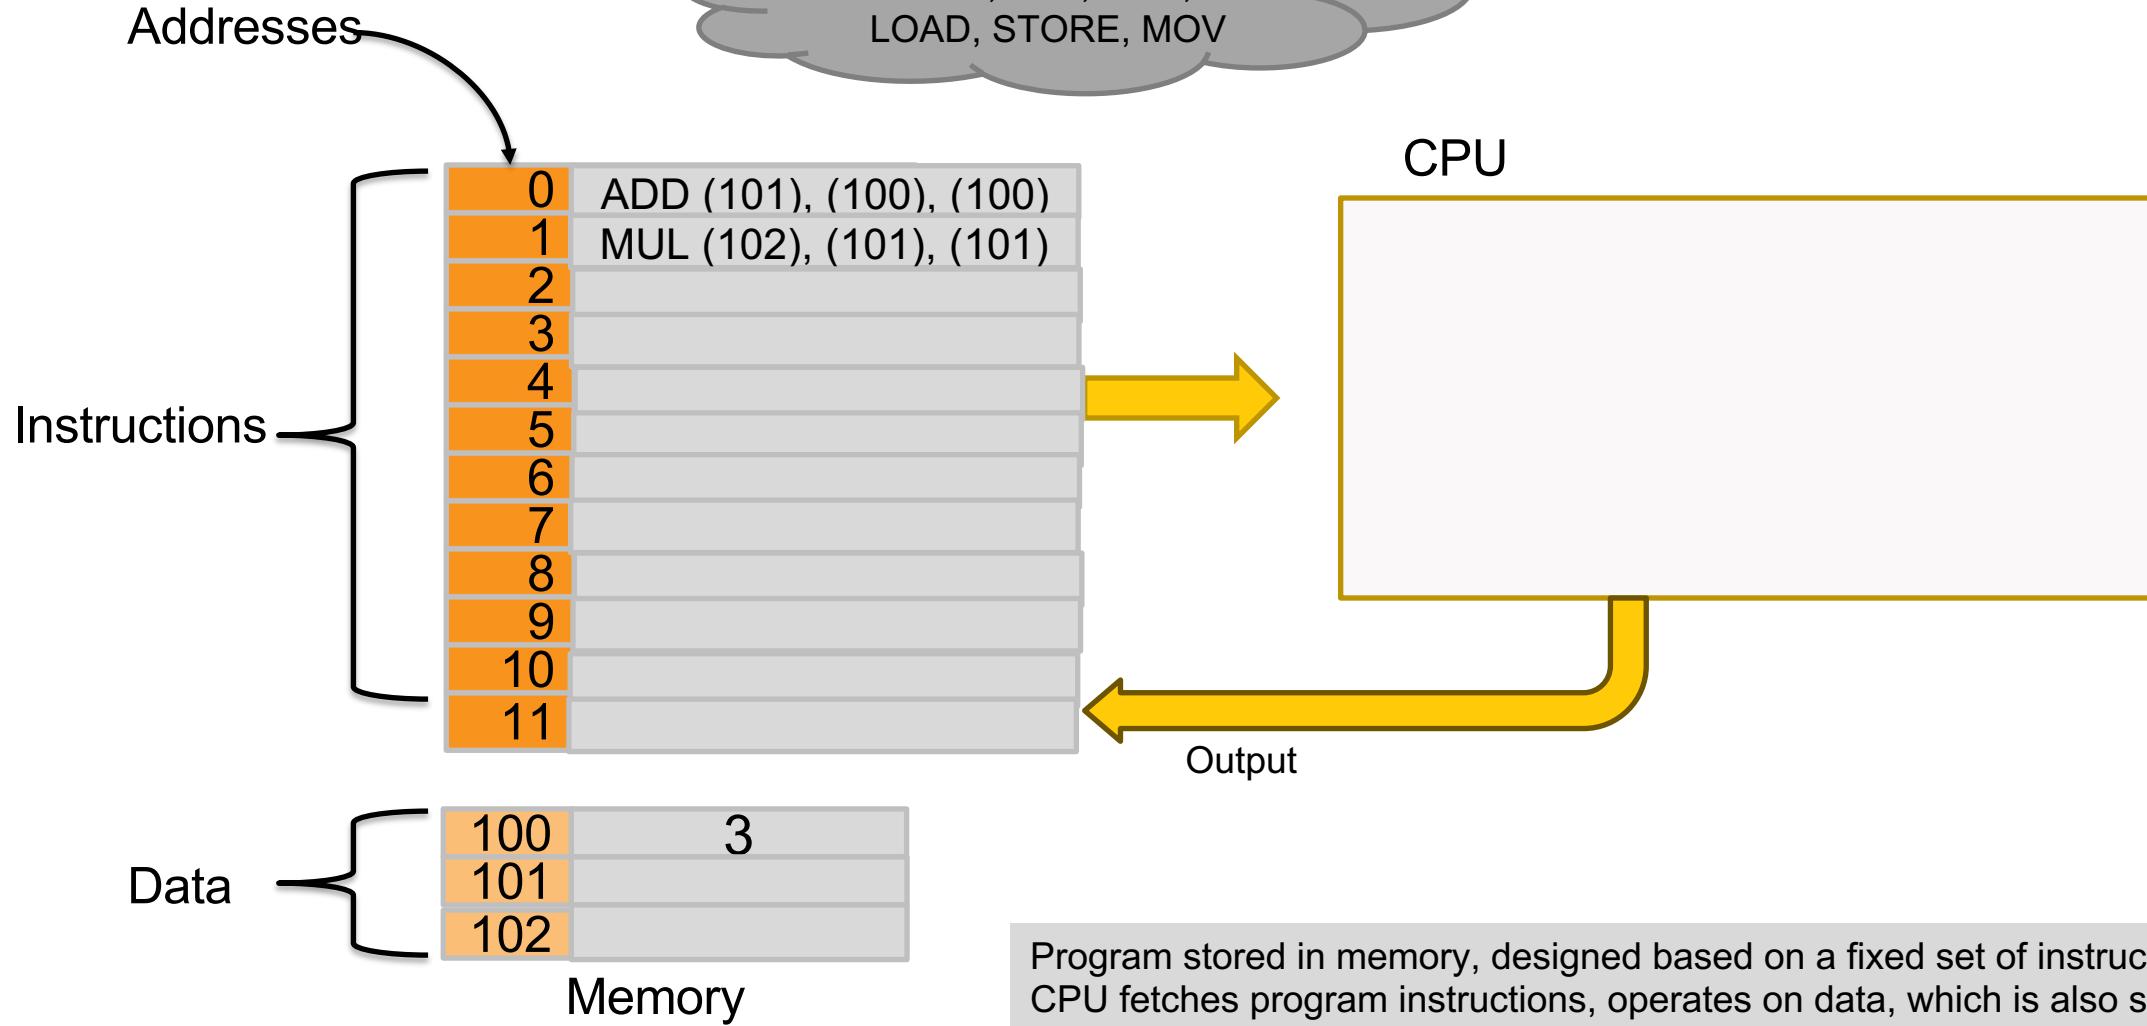
Indian Institute of Technology Madras

chester@cse.iitm.ac.in

Stored Program Concept

Instruction Set

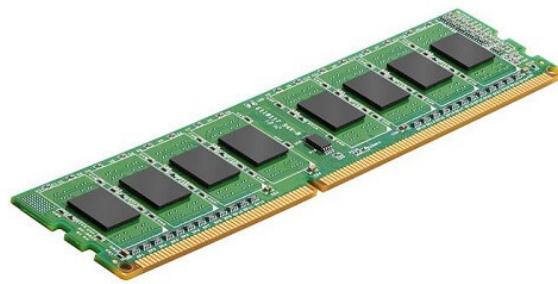
ADD, SUB, MUL, JUMP,
CMP, JEQ, NOP,
LOAD, STORE, MOV



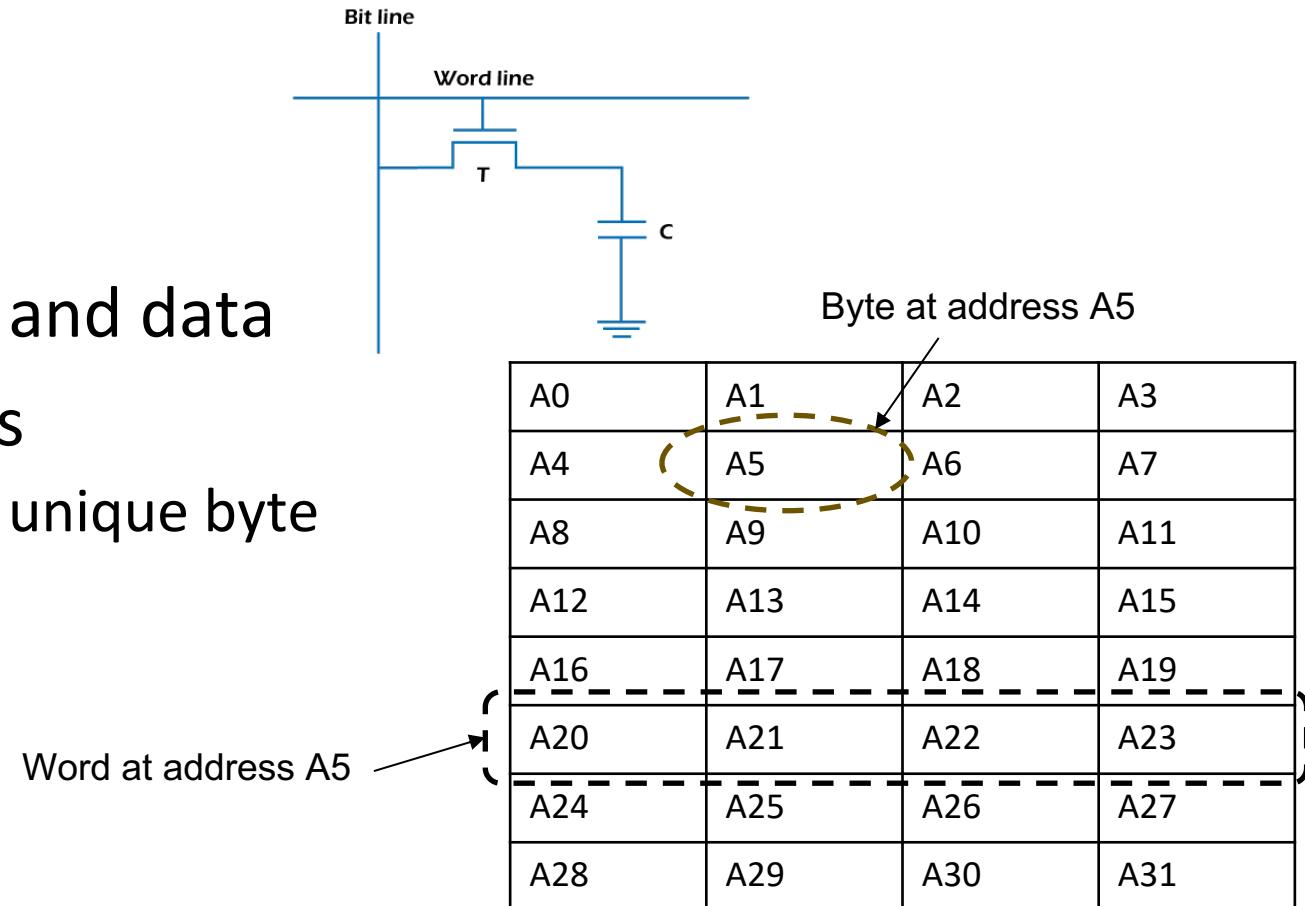
Memory (RAM)

Large

- Varies from 1MB to 1TB
- Made with capacitors
- Used to store programs and data
- Accessed with addresses
 - Each address accesses a unique byte
 - 4 bytes is called a word

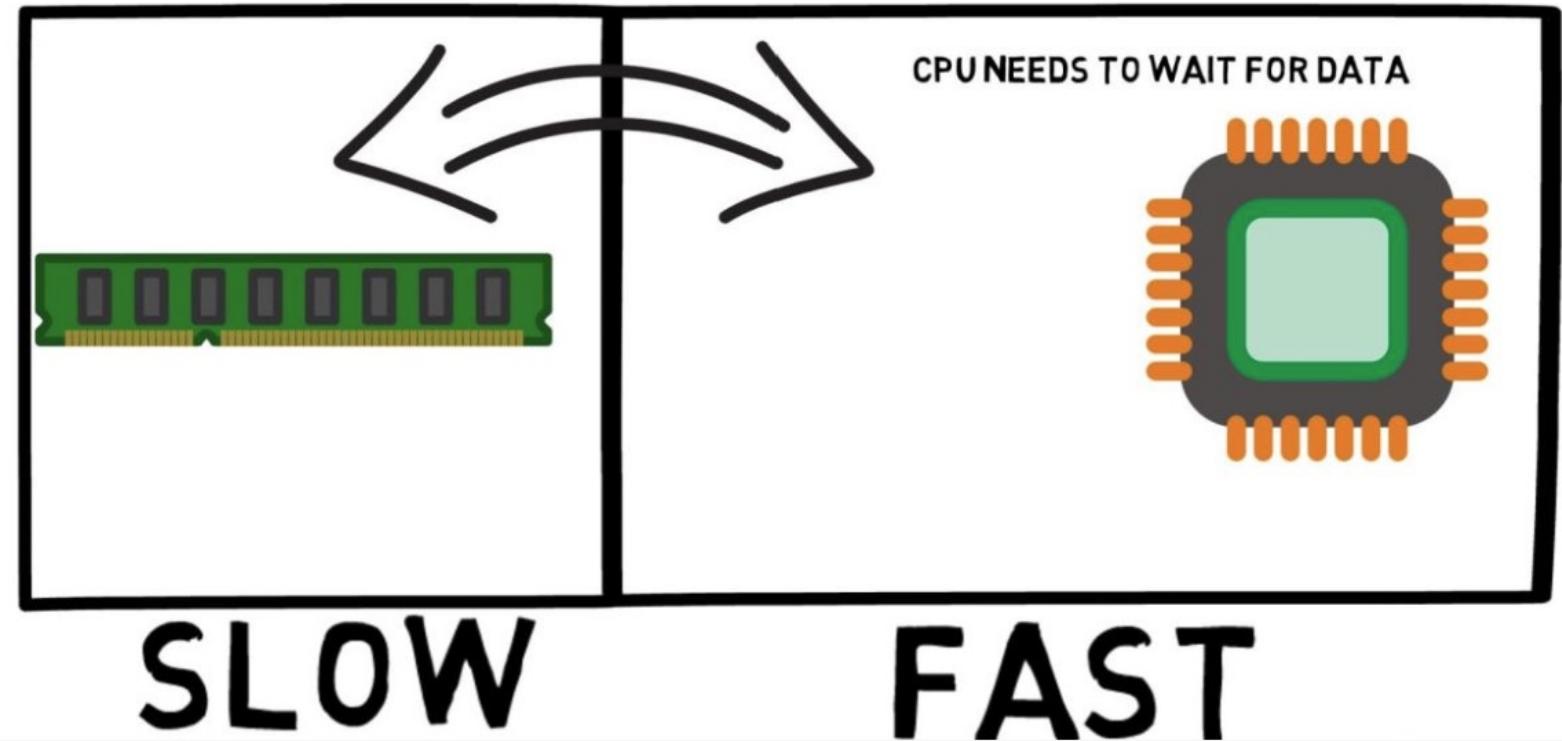


Slow



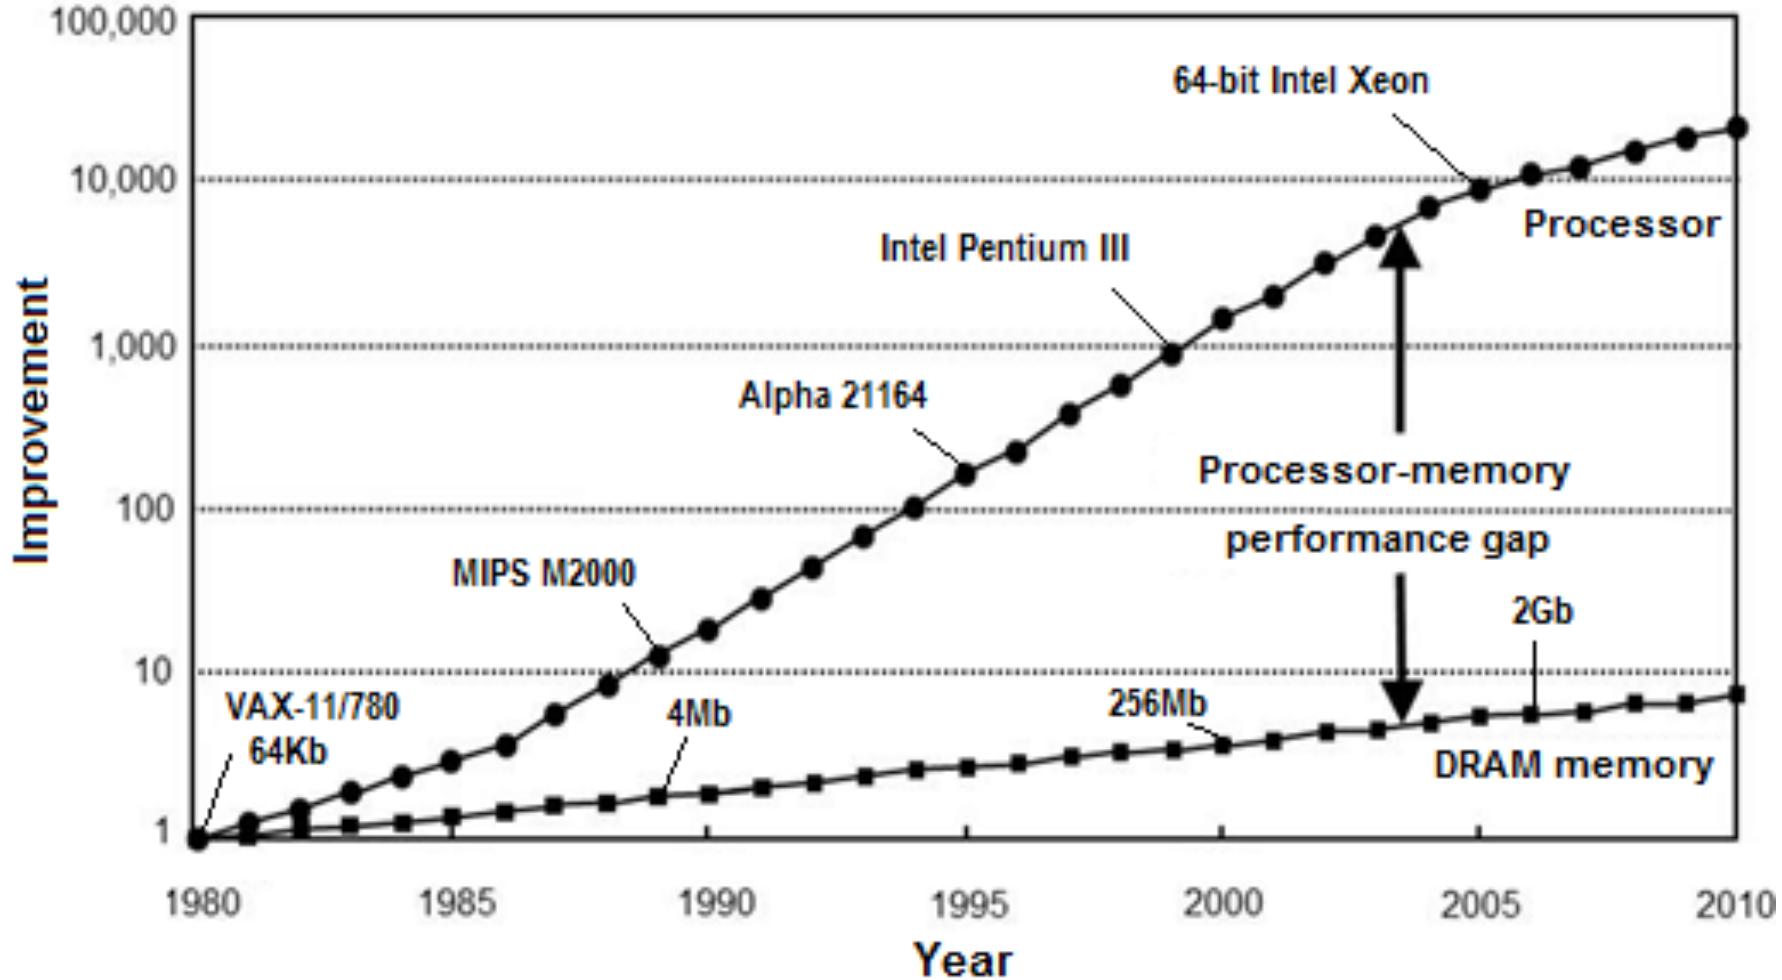
VonNeumann Bottleneck

THE VON NEUMANN BOTTLENECK



CPU too fast compared to memory. CPU

Processor-Memory Gap

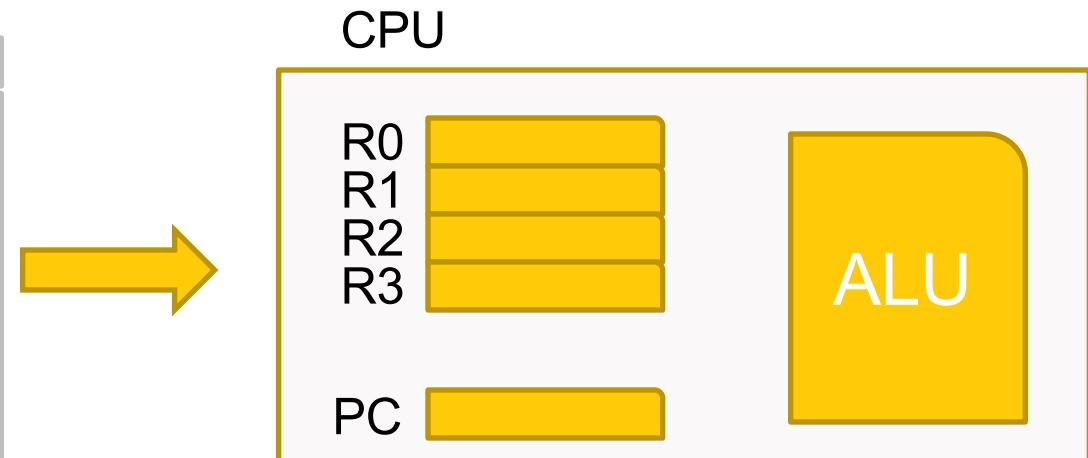
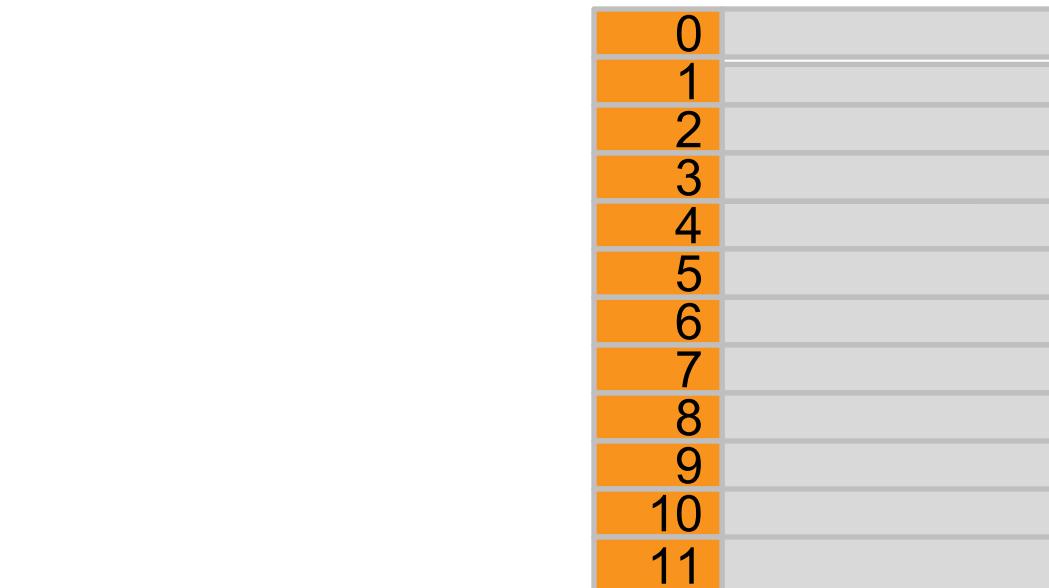


The gap has increased steadily

Registers

Instruction Set

ADD, SUB, MUL, JUMP,
CMP, JEQ, NOP,



Output



Memory

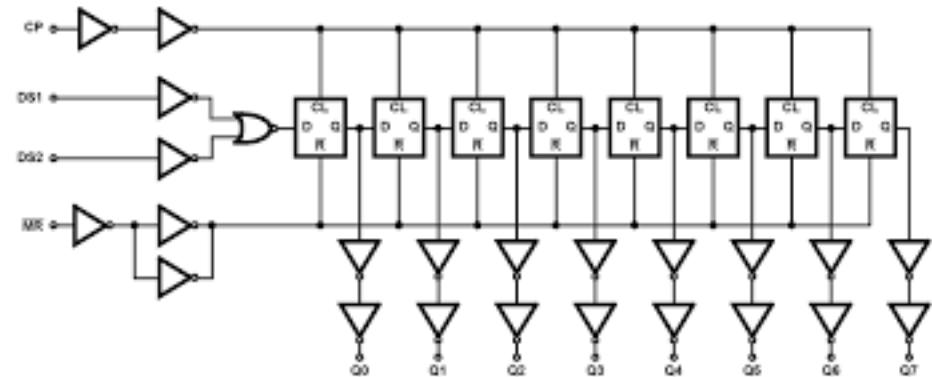
Introduce registers, which are closer to the CPU and much faster than RAM But limited

Registers

Few

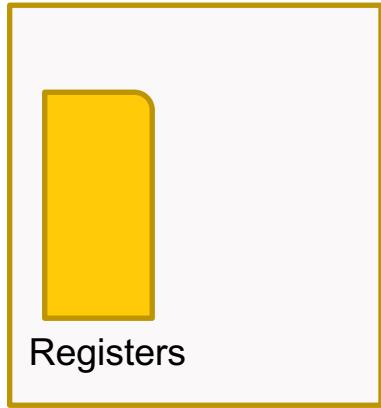
- Typically around 16, 32, or 64 registers
 - How to select the number of register?
 - Fewer is faster (easier to decode, shorter wires between connections)
 - Too few is slower 😞 (need to go to memory to access data. Can't cache much)
 - Need to find the right number
- Placed very close to the processor.
 - Operate as fast as the processor.
 - Made with CMOS transistors

Fast

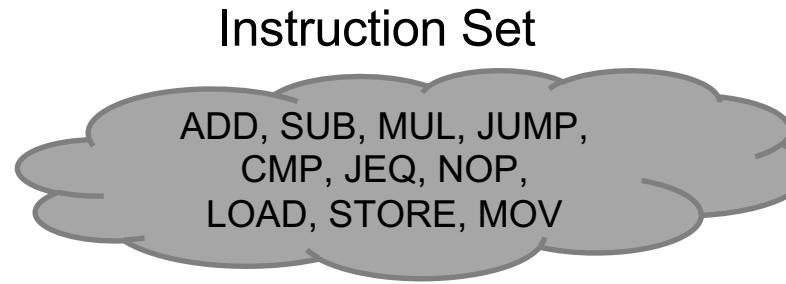


DRAMs vs Registers

Processor Core



Load / Store Instructions



0	LOAD R0, (100)
1	LOAD R1, (101)
2	ADD R0, R0
3	MUL R1, R0, R0
4	STORE R1, (102)
5	
6	
7	
8	
9	
10	
11	

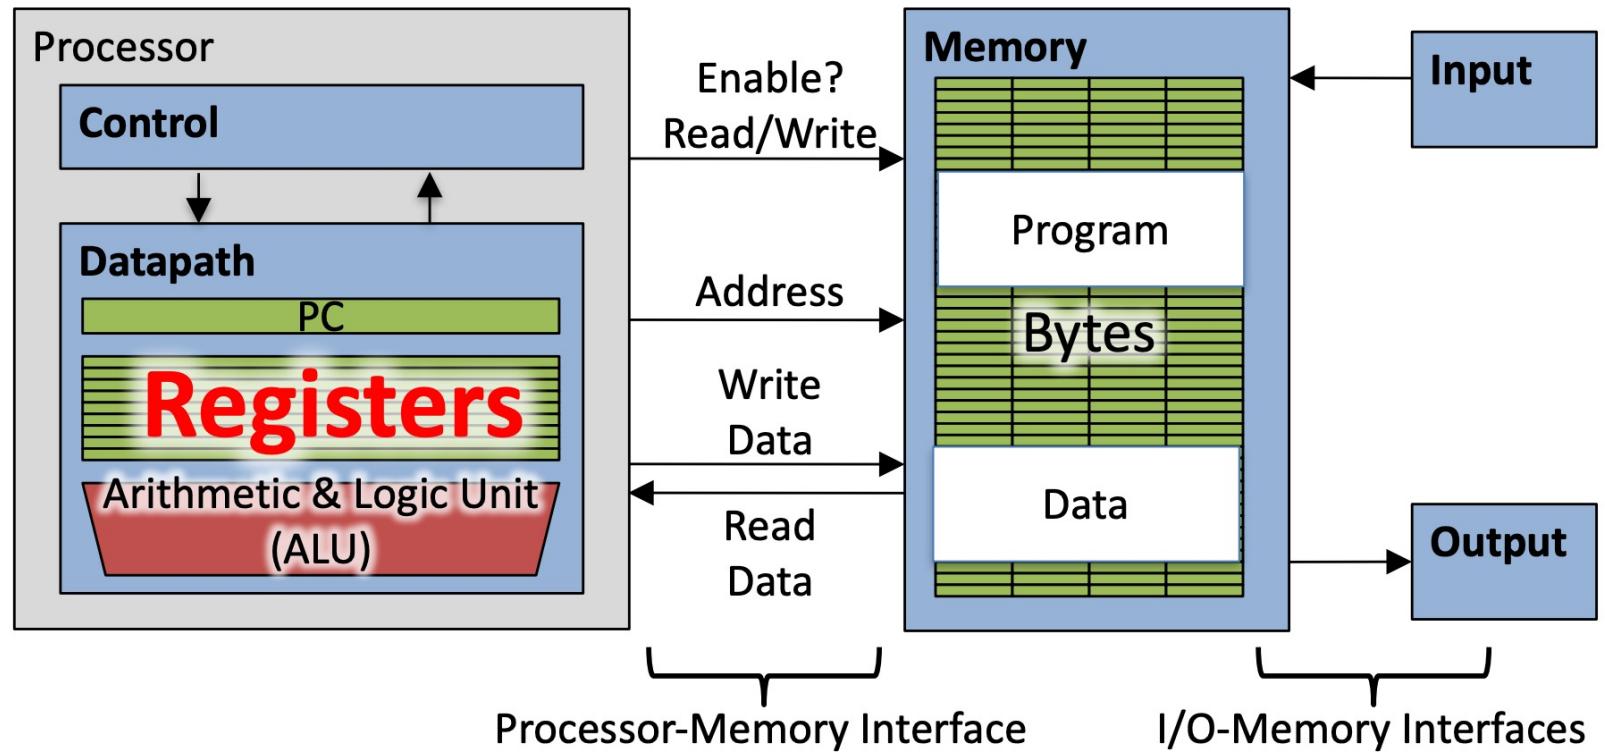


100	5
101	6
102	

Memory

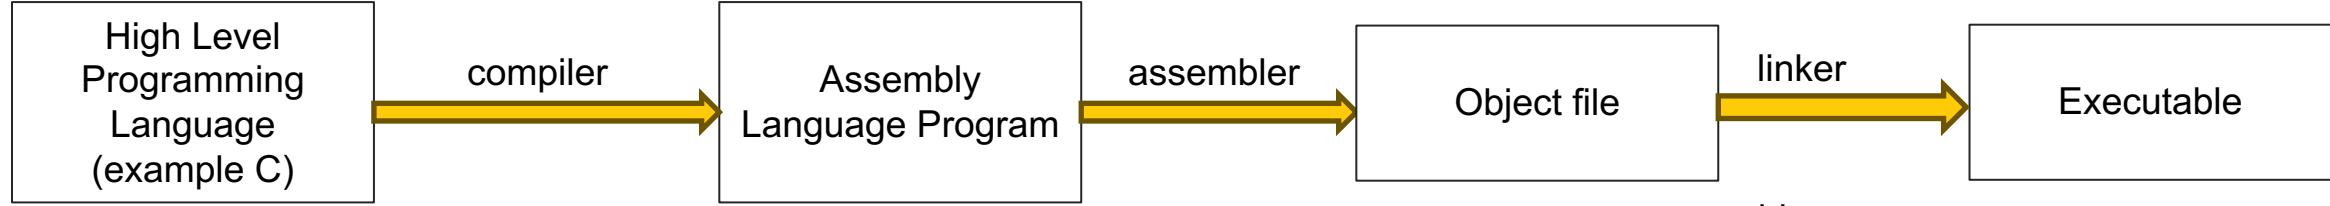
LOAD and STORE instructions to move data between memory and registers. All data now needs to be loaded first and then

Interfaces



A slightly deeper look inside the processor
Notice the Processor-Memory Interface and I/O Interfaces

Program transformations



```

3 void increment_elements(int *array, int n){
4     int i;
5     for(i=0; i<n; ++i){
6         array[i] = array[i] + 1;
7     }
8 }
9
10 void exor_elements(int *array, int n){
11    int i;
12    for(i=0; i<n; ++i){
13        array[i] = array[i] ^ 0xFFFFFFFF ;
14    }
15 }
16
17 int main(){
18     int A[10];
19     int i;
20     int choice;
21
22     for(i=0; i<10; ++i) scanf("%d", &A[i]);
23     if (A[0] == 0 && A[1]==0){
24         increment_elements(&A[2], 8);
25     }
26     else{
27         exor_elements(&A[2], 8);
28     }
29 }
  
```

C Code

```

addiw   s0,s0,-16
unimp
ld      s0,32(a2)
lui     t1,0xfffffa
bgeu   zero,a6,14
c.slli  zero,0xa
unimp
addi   s1,sp,32
lui    tp,0xffffe1
fld    fa2,376(sp)
ld     a3,80(a0)
0x7032
lw     a2,120(a4)
addiw  tp,tp,-5
fld    fa2,224(s0)
0x5f307032615f
fld    ft4,120(sp)
fld    fa2,224(s0)
0x5f307032645f
0x30703263
  
```

Assembly code

```

0000 797122f4 0018233c a4fc8e87 232af4fc
0010 232604fe 35a08327 c4fe8a07 033784fd
0020 ba979843 8327c4fe 8a078336 84fdb697
0030 05270127 98c38327 c4fe8527 2326f4fe
0040 0327c4fe 832744fd 01278127 e345f7fc
0050 01000100 22744561 82807971 22f40018
0060 233ca4fc ae87232a f4fc2326 04fe05a8
0070 8327c4fe 8a070337 84fdb97 94438327
0080 c4fe8a07 033784fd ba973687 1347f7ff
0090 012798c3 8327c4fe 85272326 f4fe0327
00a0 c4fe8327 44fd0127 8127e343 f7fc0100
00b0 01002274 45618280 397106fc 22f88000
00c0 232604fe 2da01307 04fc8327 c4fe8a07
00d0 ba97be85 b7070000 138507000 97000000
00e0 e7800000 8327c4fe 85272326 f4fe8327
00f0 c4fe1b87 0700a547 e3d7e7fc 832704fc
0100 91ef8327 44fc99eb 930704fc a107a145
0110 3e859700 0000e780 000011a8 930704fc
0120 a107a145 3e859700 0000e780 00008147
  
```

Object file
(relocatable. Addresses start from 0)

```

100b0 93070000 99c71735 00001305 e5aa6f20
100c0 30008280 97410100 9381c1ed 1385a176
100d0 17460100 1306066d 098e8145 ef000023
100e0 17250000 1305057e 19c51735 00001305
100f0 a5a7ef20 e07cef00 c01a0245 2c000146
10100 ef000e010 41a24111 22e01384 017a8347
10110 040006e4 99ef9307 000089cb 17350100
10120 1305c53e 97000000 e7000000 85472300
10130 f400a260 02644101 82809307 000099cb
10140 9385817a 17350100 1305453c 17030000
10150 67000000 82807971 22f40018 233ca4fc
10160 ae87232a f4fc2326 04fe35a0 8327c4fe
10170 8a070337 84fdb97 98438327 c4fe8a07
10180 833684fd b6970527 012798c3 8327c4fe
10190 85272326 f4fe0327 c4fe8327 44fd0127
101a0 8127e345 f7fc0100 01002274 45618280
101b0 797122f4 0018233c a4fc8e87 232af4fc
101c0 232604fe 05a88327 c4fe8a07 033784fd
101d0 ba979443 8327c4fe 8a070337 84fdb97
  
```

Executable
(addresses resolved)

Encoded instructions

Program Execution

3

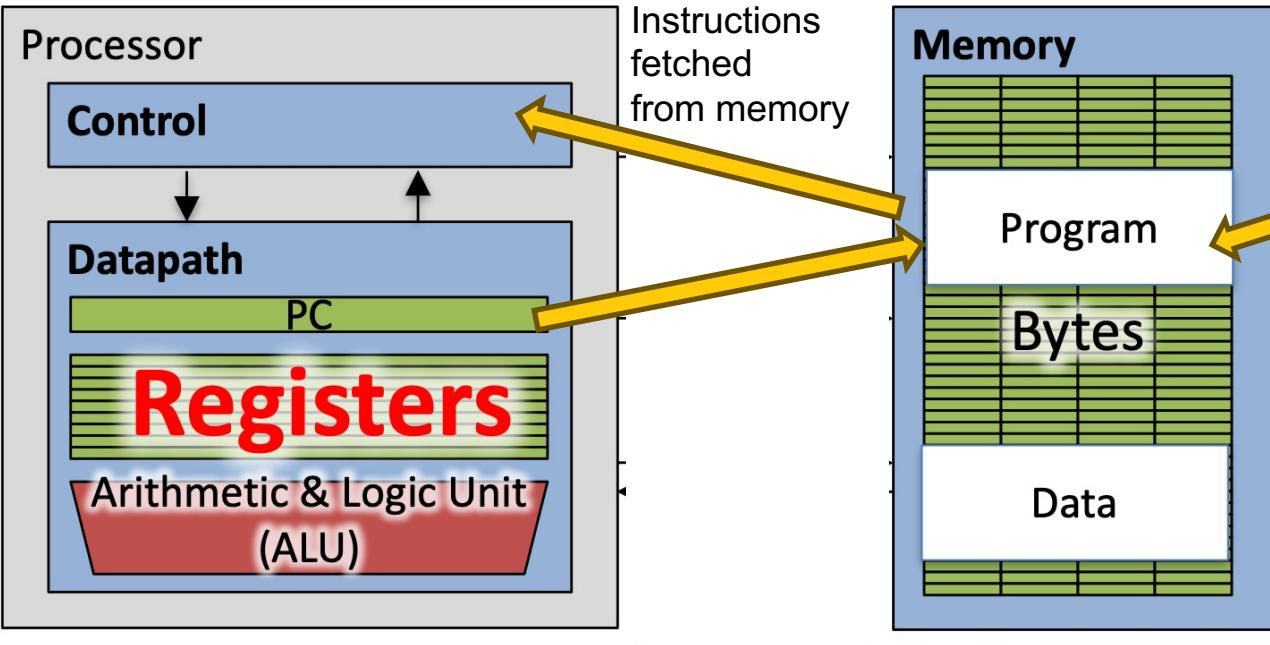
Instructions decoded
and executed

2

Instructions
fetched
from memory

1

Executable loaded
into memory
by a loader



100b0	93070000	99c71735	00001305	e5aa6f20
100c0	30008280	97410100	9381c1ed	1385a176
100d0	17460100	1306066d	098e8145	ef000023
100e0	17250000	1305057e	19c51735	00001305
100f0	a5a7ef20	e07cef00	c01a0245	2c000146
10100	ef00e010	41a24111	22e01384	017a8347
10110	04000644	99ef9307	000089cb	17350100
10120	1305c53e	97000000	e7000000	85472300
10130	f400a260	02644101	82809307	000099cb
10140	9385817a	17350100	1305453c	17030000
10150	67000000	82807971	22f40018	233ca4fc
10160	ae87232a	f4fc2326	04fe35a0	8327c4fe
10170	8a070337	84fdbba97	98438327	c4fe8a07
10180	833684fd	b6970527	012798c3	8327c4fe
10190	85272326	f4fe0327	c4fe8327	44fd0127
101a0	8127e345	f7fc0100	01002274	45618280
101b0	797122f4	0018233c	a4fciae87	232af4fc
101c0	232604fe	05a88327	c4fe8a07	033784fd
101d0	ba979443	8327c4fe	8a070337	84fdbba97

Executable

Instructions in the Program

Program Memory

100b0	93070000	99c71735	00001305	e5aa6f20
100c0	3e800093	07410100	9381c1ed	1385a176
100d0	17460100	206066d	098e8145	ef000023
100e0	17250000	1357057e	19c51735	00001305
100f0	a5a7ef20	e078af00	c01a0245	2c000146
10100	ef00e010	41a24111	22e01384	017a8347
10110	040006e4	99ef938	000089cb	17350100
10120	1305c53e	9700000	e7000000	85472300
10130	f400a260	99ef938	000089cb	17350100
10140	9385817a	10000000	00000000	00000000
10150	67000000	82807971	22f40018	233ca4fc
10160	ae87232a	f4fc2326	04fe35a0	8327c4fe
10170	8a070337	84fdb97	98438327	c4fe8a07
10180	833684fd	b6970527	012798c3	8327c4fe
10190	85272326	f4fe0327	c4fe8327	44fd0127
101a0	8127e345	f7fc0100	01002274	45618280
101b0	797122f4	0018233c	a4fcae87	232af4fc
101c0	232604fe	05a88327	c4fe8a07	033784fd
101d0	ba979443	8327c4fe	8a070337	84fdb97

addi x1 , x0, 1000



- Each 32-bit in the program memory corresponds to an encoded instruction
- Unique map from instruction to encoding
- Instructions picked from the processor's Instruction Set Architecture



The Instruction Set

The language that CPUs understand

Different Instruction Set: x86, ARM, **RISC V**, MIPS.

- Program written for one processor cannot execute on another
- Early trend: add more instructions to new CPUs for elaborate operations
- RISC— Reduced Instruction Set Computing –
 - Instructions small and simple
 - Software does complicated operations by composing simpler ones

Any instruction set works over two aspects.

- Set of operations (opcodes)
- Set of registers

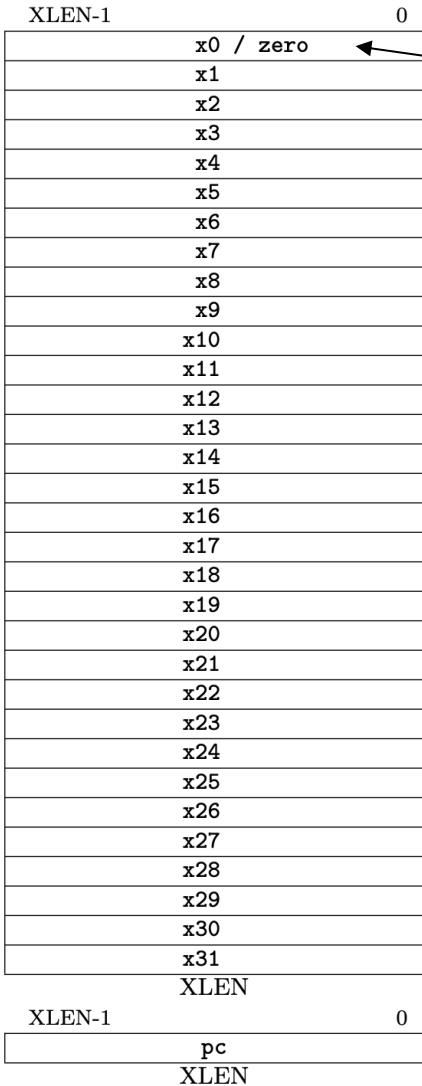


RISC V Instruction Set

- Fifth generation of RISC design from UC Berkeley
- License-free, royalty-free RISC ISA specification
- Experiencing rapid uptake in both industry and academia
- Both proprietary and open-source core implementations
- Appropriate for all levels of computing system, from microcontrollers to supercomputers
- Standard maintained by non-profit RISC-V Foundation

RISC V General Purpose Registers

XLEN is 32 or 64 bit, corresponding to 32 and 64 bit processors respectively

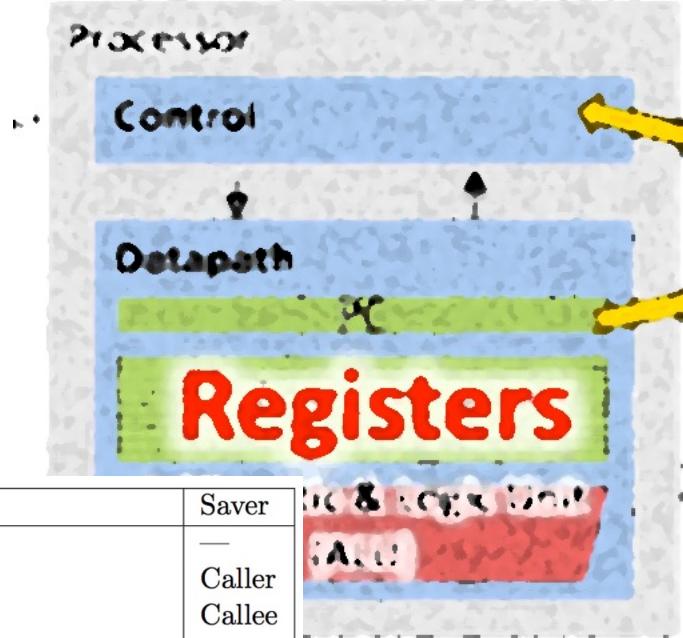


Holds zero

Numbers hardware understands

Human-friendly symbolic names in assembly code

Register	ABI Name	Description	Saver
x0	zero	Hard-wired zero	—
x1	ra	Return address	Caller
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	—
x4	tp	Thread pointer	—
x5	t0	Temporary/alternate link register	Caller
x6-7	t1-2	Temporaries	Caller
x8	s0/fp	Saved register/frame pointer	Callee
x9	s1	Saved register	Callee
x10-11	a0-1	Function arguments/return values	Caller
x12-17	a2-7	Function arguments	Caller
x18-27	s2-11	Saved registers	Callee
x28-31	t3-6	Temporaries	Caller

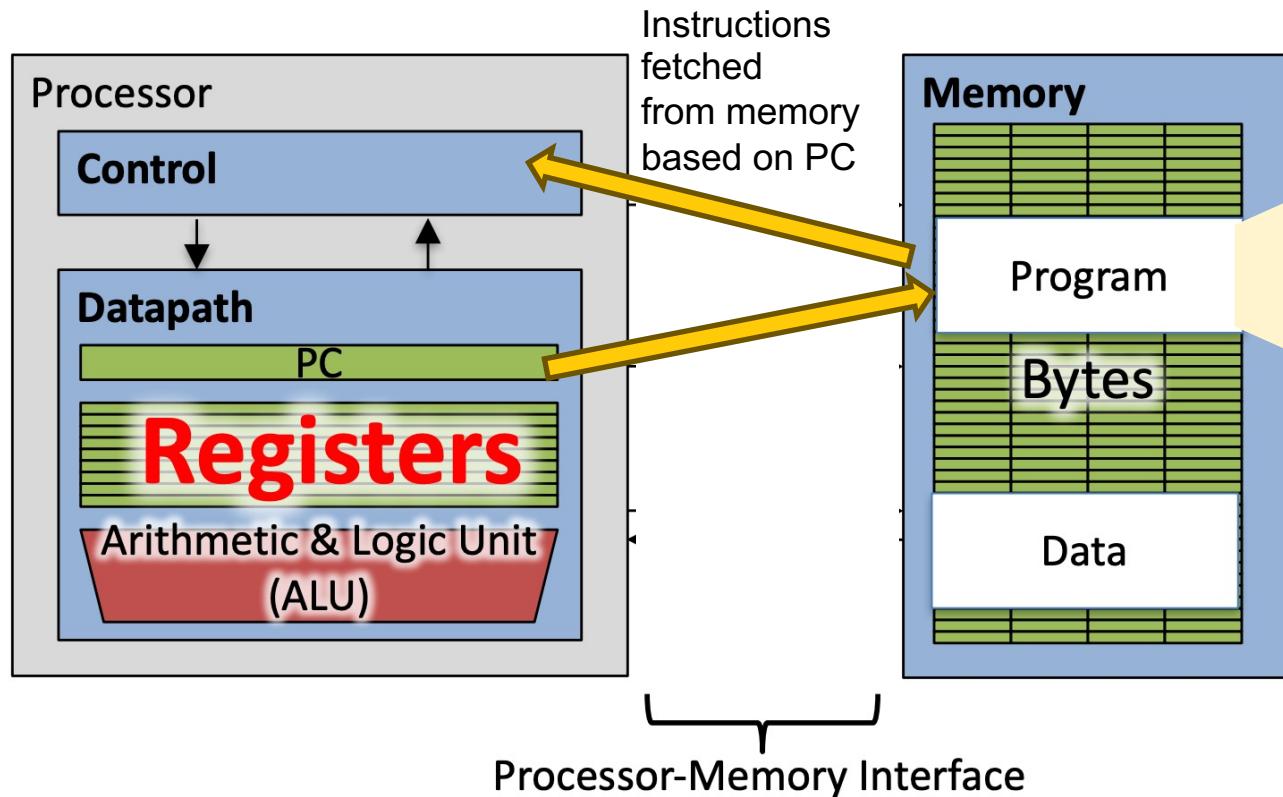


Processor

The Program Counter (pc)

PC: Points to the current instruction being executed

Typically increments in 4 bytes.



Instruction fetched when
PC = 0x100c4

100b0	93070000	99c71735	00001305	e5aa6f20
100c0	30008280	97410100	9381c1ed	1385a176
100d0	17460100	1306066d	098e8145	ef000023
100e0	17250000	1305057e	19c51735	00001305
100f0	a5a7ef20	e07cef00	c01a0245	2c000146
10100	ef00e010	41a24111	22e01384	017a8347
10110	040006e4	99ef9307	000089cb	17350100
10120	1305c53e	97000000	e7000000	85472300
10130	f400a260	02644101	82809307	000099cb
10140	9385817a	17350100	1305453c	17030000
10150	67000000	82807971	22f40018	233ca4fc
10160	ae87232a	f4fc2326	04fe35a0	8327c4fe
10170	8a070337	84fdb97	98438327	c4fe8a07
10180	833684fd	b6970527	012798c3	8327c4fe
10190	85272326	f4fe0327	c4fe8327	44fd0127
101a0	8127e345	f7fc0100	01002274	45618280
101b0	797122f4	0018233c	a4fcae87	232af4fc
101c0	232604fe	05a88327	c4fe8a07	033784fd
101d0	ba979443	8327c4fe	8a070337	84fdb97

Program in Memory



Programs

```
3 void increment_elements(int *array, int n){
4     int i;
5     for(i=0; i<n; ++i){
6         array[i] = array[i] + 1;
7     }
8 }
9
10 void exor_elements(int *array, int n){
11     int i;
12     for(i=0; i<n; ++i){
13         array[i] = array[i] ^ 0xFFFFFFFF ;
14     }
15 }
16
17 int main(){
18     int A[10];
19     int i;
20     int choice;
21
22     for(i=0; i<10; ++i) scanf("%d", &A[i]);
23     if (A[0] == 0 && A[1]==0){
24         increment_elements(&A[2], 8);
25     }
26     else{
27         exor_elements(&A[2], 8);
28     }
29 }
```



Instructions for every program construct

```
3 void increment_elements(int *array, int n){ → Arithmetic operations
4     int i;
5     for(i=0; i<n; ++i){ → Reading and Writing data (Data Transfer)
6         array[i] = array[i] + 1;
7     }
8 }
9

10 void exor_elements(int *array, int n){
11     int i; → Local variables
12     for(i=0; i<n; ++i){
13         array[i] = array[i] ^ 0xFFFFFFFF ;
14     }
15 }
16

17 int main(){
18     int A[10];
19     int i;
20     int choice;
21
22     for(i=0; i<10; ++i) scanf("%d", &A[i]); → Reading from I/O (keyboard)
23     if (A[0] == 0 && A[1]==0){ → Decision making operations
24         increment_elements(&A[2], 8);
25     }
26     else{
27         exor_elements(&A[2], 8); → Functional calls
28     }
29 }
```

Typical instructions

C statement

```
a = b + c;
```

Assembly instruction with operands in registers

```
add x1, x2, x3 # x1 = x2 + x3
```

Operation code
(opcode)

Destination
register

First source
register

Second source
register

Typical instructions

C statement

```
a = b + c;
```

C statement

```
a = b - 10;
```

Assembly instruction with operands in registers

```
add x1, x2, x3 # x1 = x2 + x3
```

Operation code
(opcode)

Destination
Register (rd)

First source
register (rs1)

Second source
register (rs2)

Assembly instruction with Immediate operands

```
addi x1, x2, -10 # x1 = x2 - 10
```

Operation code
(opcode)

Destination
Register (rd)

First source
register (rs1)

Immediate operand



Instruction Encoding

add x1, x2, x3 # x1 = x2 + x3

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
funct7		rs2		rs1		funct3		rd			opcode			R-type
0000	000		00011		00010		000		00001		0110011		0x003100B3	

add x1, x3, x2 # x1 = x3 + x2

0000	000		00010		00011		000		00001		0110011		0x002180B3
------	-----	--	-------	--	-------	--	-----	--	-------	--	---------	--	------------

sub x1, x3, x2 # x1 = x3 - x2

0100	000		00010		00011		000		00001		0110011		0x402180B3
------	-----	--	-------	--	-------	--	-----	--	-------	--	---------	--	------------

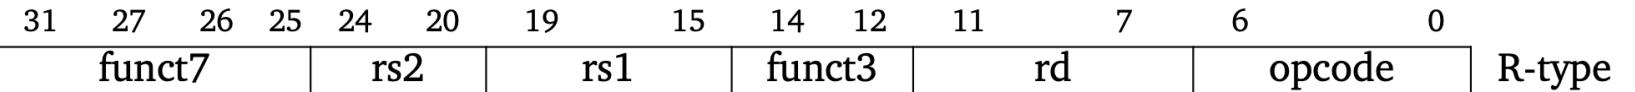
Ref.. [RISCV_CARD.PDF](#)

Try it out here.. [RISCV ONLINE ASSEMBLER](#)

R-Type Instructions in RISC V

R-Type

Inst	Name	FMT	Opcode	funct3	funct7	Description (C)	Note
add	ADD	R	0110011	0x0	0x00	$rd = rs1 + rs2$	
sub	SUB	R	0110011	0x0	0x20	$rd = rs1 - rs2$	
xor	XOR	R	0110011	0x4	0x00	$rd = rs1 \oplus rs2$	
or	OR	R	0110011	0x6	0x00	$rd = rs1 rs2$	
and	AND	R	0110011	0x7	0x00	$rd = rs1 \& rs2$	
sll	Shift Left Logical	R	0110011	0x1	0x00	$rd = rs1 << rs2$	
srl	Shift Right Logical	R	0110011	0x5	0x00	$rd = rs1 >> rs2$	
sra	Shift Right Arith*	R	0110011	0x5	0x20	$rd = rs1 >> rs2$	msb-extends
slt	Set Less Than	R	0110011	0x2	0x00	$rd = (rs1 < rs2)?1:0$	
sltu	Set Less Than (U)	R	0110011	0x3	0x00	$rd = (rs1 < rs2)?1:0$	zero-extends Unsigned comparison



Shift Left Logical (shift left x12 by 2-bits)

```
sll x12, x12, x11      # x12 = x12 << x1
```

Initial values: x11 = 0x2, x12 = 0x1

Final values: x12=0x4

Shift Right Logical (shift right x12 by 2-bits)

```
srl x12, x12, x11      # x12 = x12 >> x11
```

Initial values: x11 = 0x2, x12 = 0xabcd5043

Final values: x12=0x2af35410

Shift Right Arithmetic (shift right x12 by 2-bits) and extend msb

```
sra x12, x12, x11      # x12 = x12 >> x1
```

Initial values: x11 = 0x2, x12 = 0xabcd5043

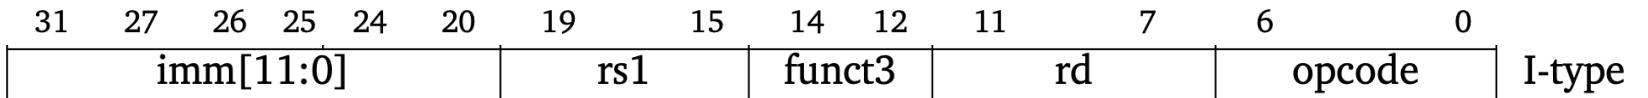
Final values: x12=0xeaf35410

Try them out here: [RISC V Online Interpreter](#)

I-Type Instructions in RISC V

I-Type

Inst	Name	FMT	Opcode	funct3	funct7	Description (C)	Note
addi	ADD Immediate	I	0010011	0x0		$rd = rs1 + imm$	
xori	XOR Immediate	I	0010011	0x4		$rd = rs1 \wedge imm$	
ori	OR Immediate	I	0010011	0x6		$rd = rs1 \vee imm$	
andi	AND Immediate	I	0010011	0x7		$rd = rs1 \wedge imm$	
slli	Shift Left Logical Imm	I	0010011	0x1	imm[5:11]=0x00	$rd = rs1 \ll imm[0:4]$	max shift is 16
srli	Shift Right Logical Imm	I	0010011	0x5	imm[5:11]=0x00	$rd = rs1 \gg imm[0:4]$	
srai	Shift Right Arith Imm	I	0010011	0x5	imm[5:11]=0x20	$rd = rs1 \gg imm[0:4]$	msb-extends
slti	Set Less Than Imm	I	0010011	0x2		$rd = (rs1 < imm)?1:0$	
sltiu	Set Less Than Imm (U)	I	0010011	0x3		$rd = (rs1 < imm)?1:0$	zero-extends

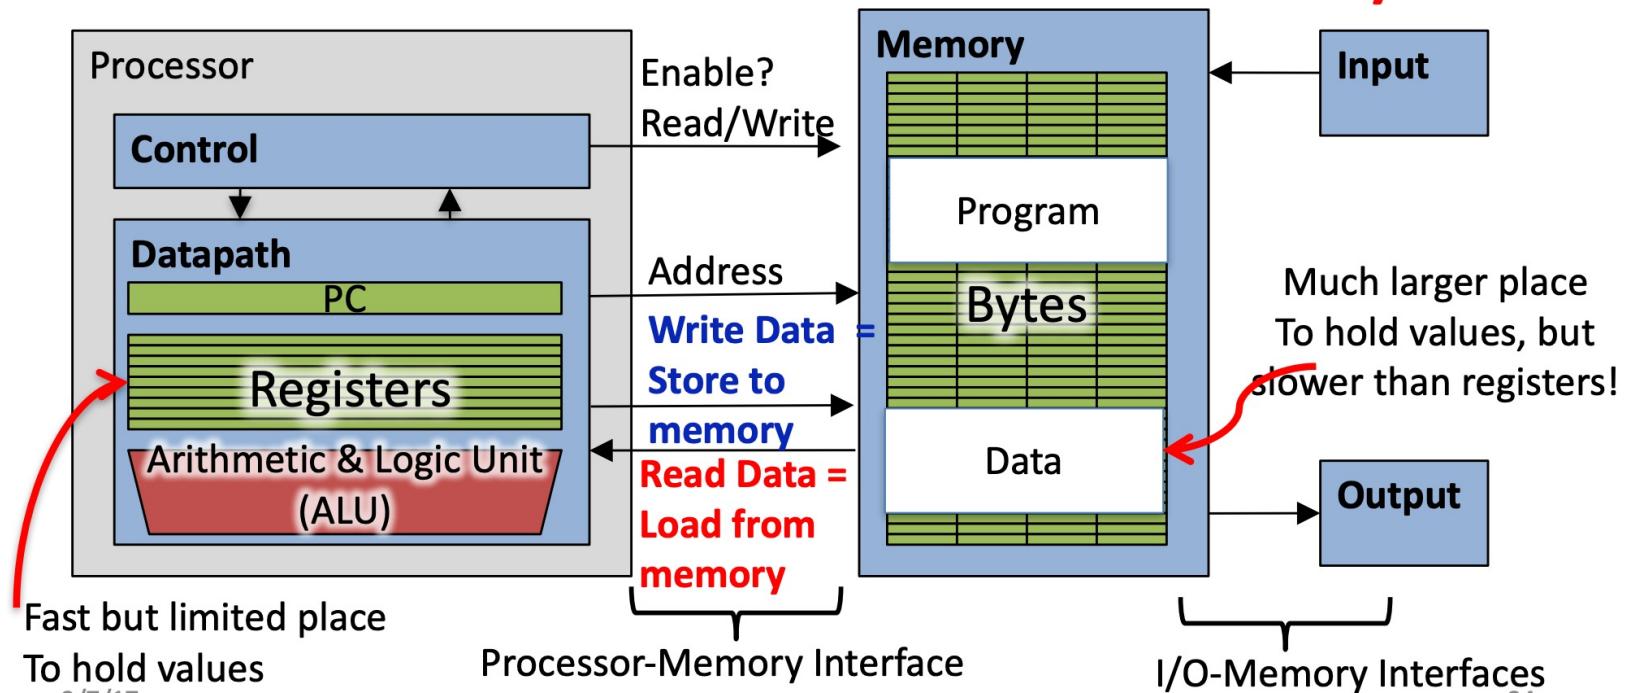


Immediate is sign extended and is of 12 bits

```
addi x12, x12, 2047 # x12 = x12 + 0x000007FF
addi x13, x13, 2048 # x13 = x13 + 0xFFFFF800
```

Initial values: x12 = x13 = 0x1
 Final values: x12 = 2048
 x13 = -2047

Memory Load and Store Operations



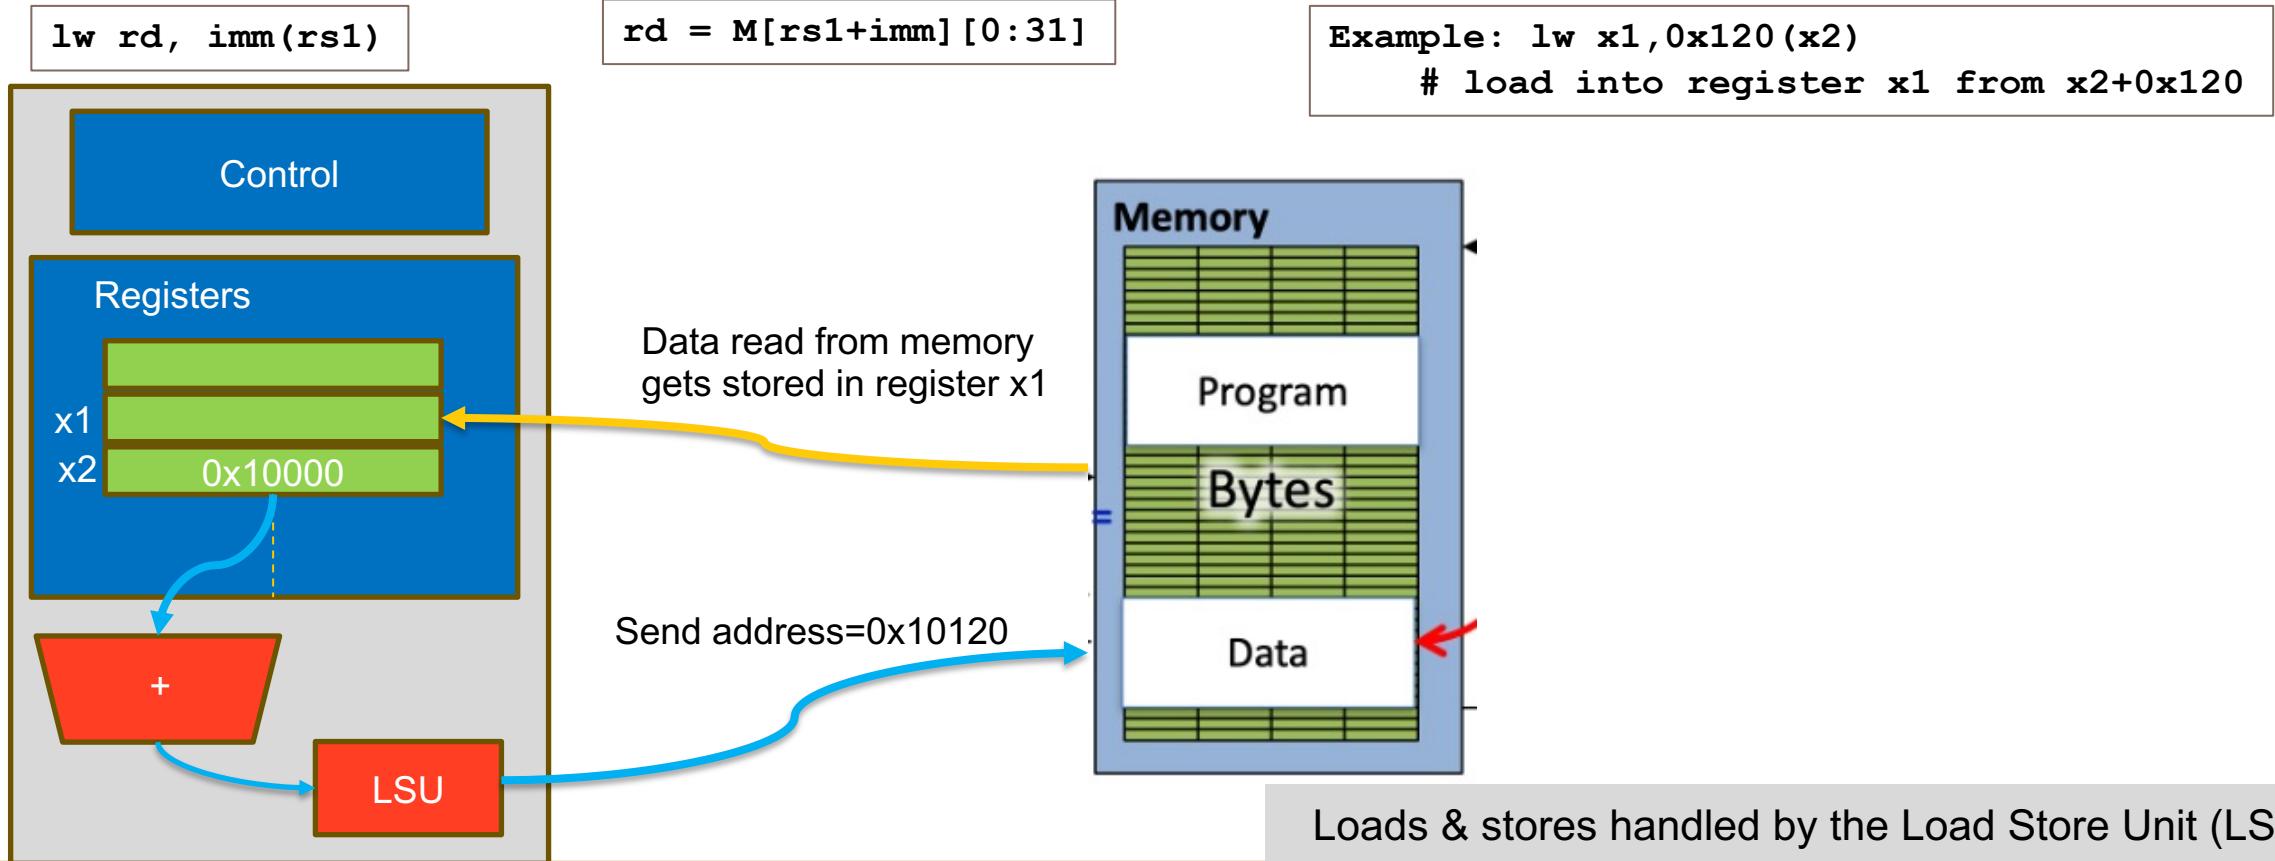
Load : Read data from memory at the given address to a register
 Store: Write data from a register to memory at the given address

Instructions for Data Transfer

data transfer instruction

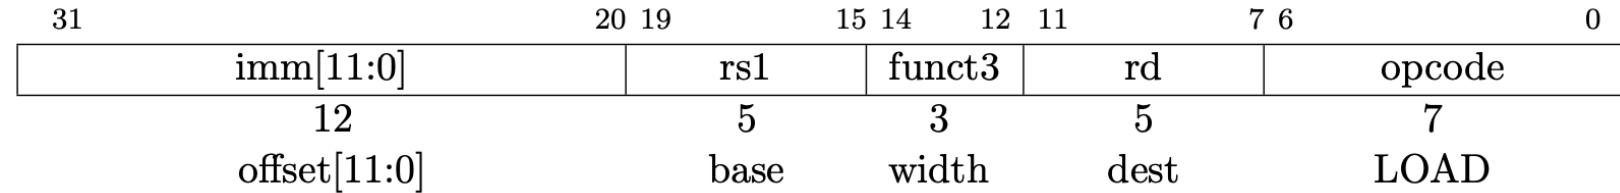
A command that moves data between memory and registers.

Load 32 bits into destination register rd from address rs1+imm; rs is the source register, imm a 12-bit signed immediate



Load Instruction Encoding

`lw rd, imm(rs1)`



- Imm is 12-bits and signed. Thus range of immediate is -2048 to +2047

C construct: arrays

```
int a[10];
int b = a[5];
```

dest: rd

base register: rs1

offset: imm

Equivalent assembly instruction

```
lw x10, 24(x12)
```

dest: rd

base of array

offset: imm



Instructions for Data Transfer

data transfer instruction

A command that moves data between memory and registers.

Load 32 bits into destination register rd from address rs+imm; rs is the source register, imm a 12-bit signed immediate

`lw rd, imm(rs)`

$rd = M[rs+imm][0:31]$

Example: `lw x1,128(x2)`



Instructions for Data Transfer

data transfer instruction

A command that moves data between memory and registers.

Load 32 bits into destination register rd from address rs+imm; rs is the source register, imm a 12-bit signed immediate

lw rd, imm(rs)

rd = M[rs+imm] [0:31]

Example: lw x1,128(x2)

Load 8 bits into destination register rd from address rs+imm; rs is the source register, imm a 12-bit signed immediate
rs is sign extended

lb rd, imm(rs)

rd = M[rs+imm] [0:7]

Example: lw x1,128(x2)

Load 8 bits into destination register rd from address rs+imm; rs is the source register, imm a 12-bit signed immediate
rs is zero extended

lbu rd, imm(rs)

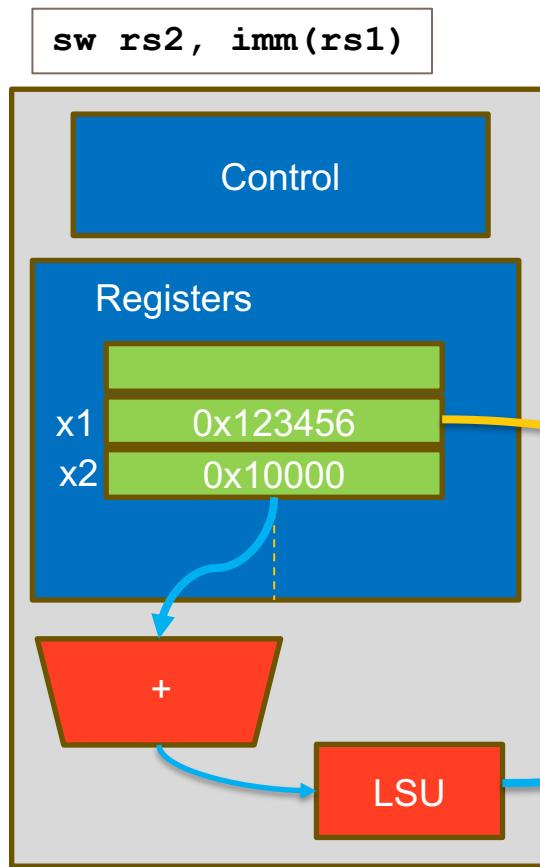
rd = M[rs+imm] [0:7]

Example: lbu x1,128(x2)

Similarly, **lh** and **lhu** are 16-bit loads that are equivalent to **lb** and **lbu**

Store Instruction

Store (write) 32 bits into destination register rd from address rs1+imm; rs is the source register, imm a 12-bit signed immediate

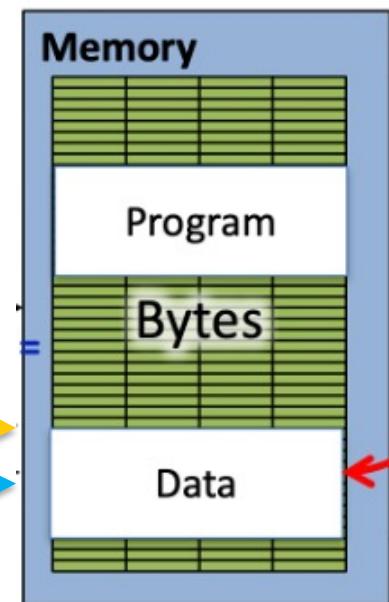


$M[rs1+imm][0:31] = rs2$

Example: sw x1,0x120(x2)
 # store into location x2 + 0x120
 the contents of x1

Data 0x123456 written to
 Memory at address 0x10120

Send address=0x10120



Store Instruction Encoding

sw rs2, imm(rs1)

31	25 24	20 19	15 14	12 11	7 6	0
imm[11:5]	rs2	rs1	funct3	imm[4:0]	opcode	
7 offset[11:5]	5 src	5 base	3 width	5 offset[4:0]	7 STORE	

- Imm is 12-bits and signed. Thus, range of immediate is -2048 to +2047

C construct: arrays

```
int a[10];
a[6] = b;
```

base register: rs1

offset: imm

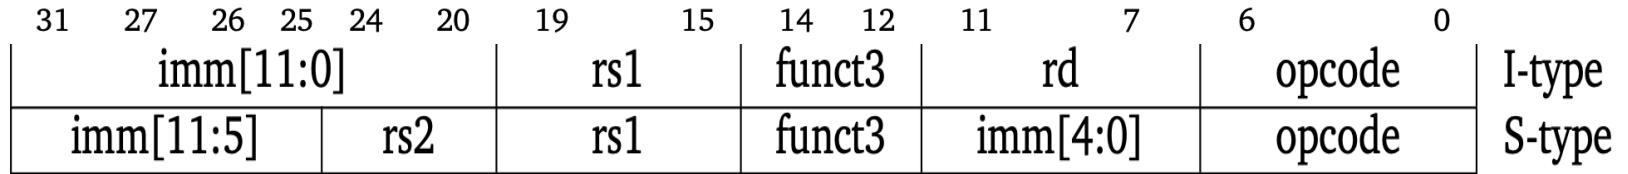
Equivalent assembly instruction

sw x10, 24(x12)

src: rs2



Other Load and Store Instructions



Inst	Name	FMT	Opcode	funct3	funct7	Description (C)	Note
lb	Load Byte	I	0000011	0x0		$rd = M[rs1+imm][0:7]$	
lh	Load Half	I	0000011	0x1		$rd = M[rs1+imm][0:15]$	
lw	Load Word	I	0000011	0x2		$rd = M[rs1+imm][0:31]$	
lbu	Load Byte (U)	I	0000011	0x4		$rd = M[rs1+imm][0:7]$	zero-extends
lhu	Load Half (U)	I	0000011	0x5		$rd = M[rs1+imm][0:15]$	zero-extends
sb	Store Byte	S	0100011	0x0		$M[rs1+imm][0:7] = rs2[0:7]$	
sh	Store Half	S	0100011	0x1		$M[rs1+imm][0:15] = rs2[0:15]$	
sw	Store Word	S	0100011	0x2		$M[rs1+imm][0:31] = rs2[0:31]$	



Example load and store

```
addi    x1, x0, 0xABCD E00  
sw      x1, 4(x2)  
lb      x3, 5(x2)  
lbu     x4, 5(x2)
```

If x1 has 0xABCD E00

After execution

x3 will hold, 0xFFFF FFEF
x4 will hold, 0x0000 00EF

Aligned vs Non-aligned memory accesses

For words

Loads and stores typically optimized for word aligned addresses. Multiple of 4 bytes.

Aligned load: From an address that is a multiple of word size (word size here is 4 bytes). Most often used.

```
x2 holds A16
```

```
lw x1, 0x0(x2)
```

Non-aligned load: From an address that is not a multiple of word size (word size here is 4 bytes). Not always supported in all architectures. Slow. May be more difficult to support due to atomicity.

```
x2 holds A6
```

```
lw x1, 0x0(x2)
```

Memory

A0	A1	A2	A3
A4	A5	A6	A7
A8	A9	A10	A11
A12	A13	A14	A15
A16	A17	A18	A19
A20	A21	A22	A23
A24	A25	A26	A27
A28	A29	A30	A31

Aligned vs Non-aligned memory accesses

For half words

Loads and stores for half words typically optimized for $\frac{1}{2}$ word aligned addresses. Multiple of 2 bytes.

Aligned: From an address that is a multiple of $\frac{1}{2}$ word size (word size here is 4 bytes). Most often used.

```
x2 holds A6
```

```
lhu x1, 0x0(x2)
```

Non-aligned : From an address that is not a multiple of word size (word size here is 4 bytes). Not always supported in all architectures. Slow. May be more difficult to support due to atomicity.

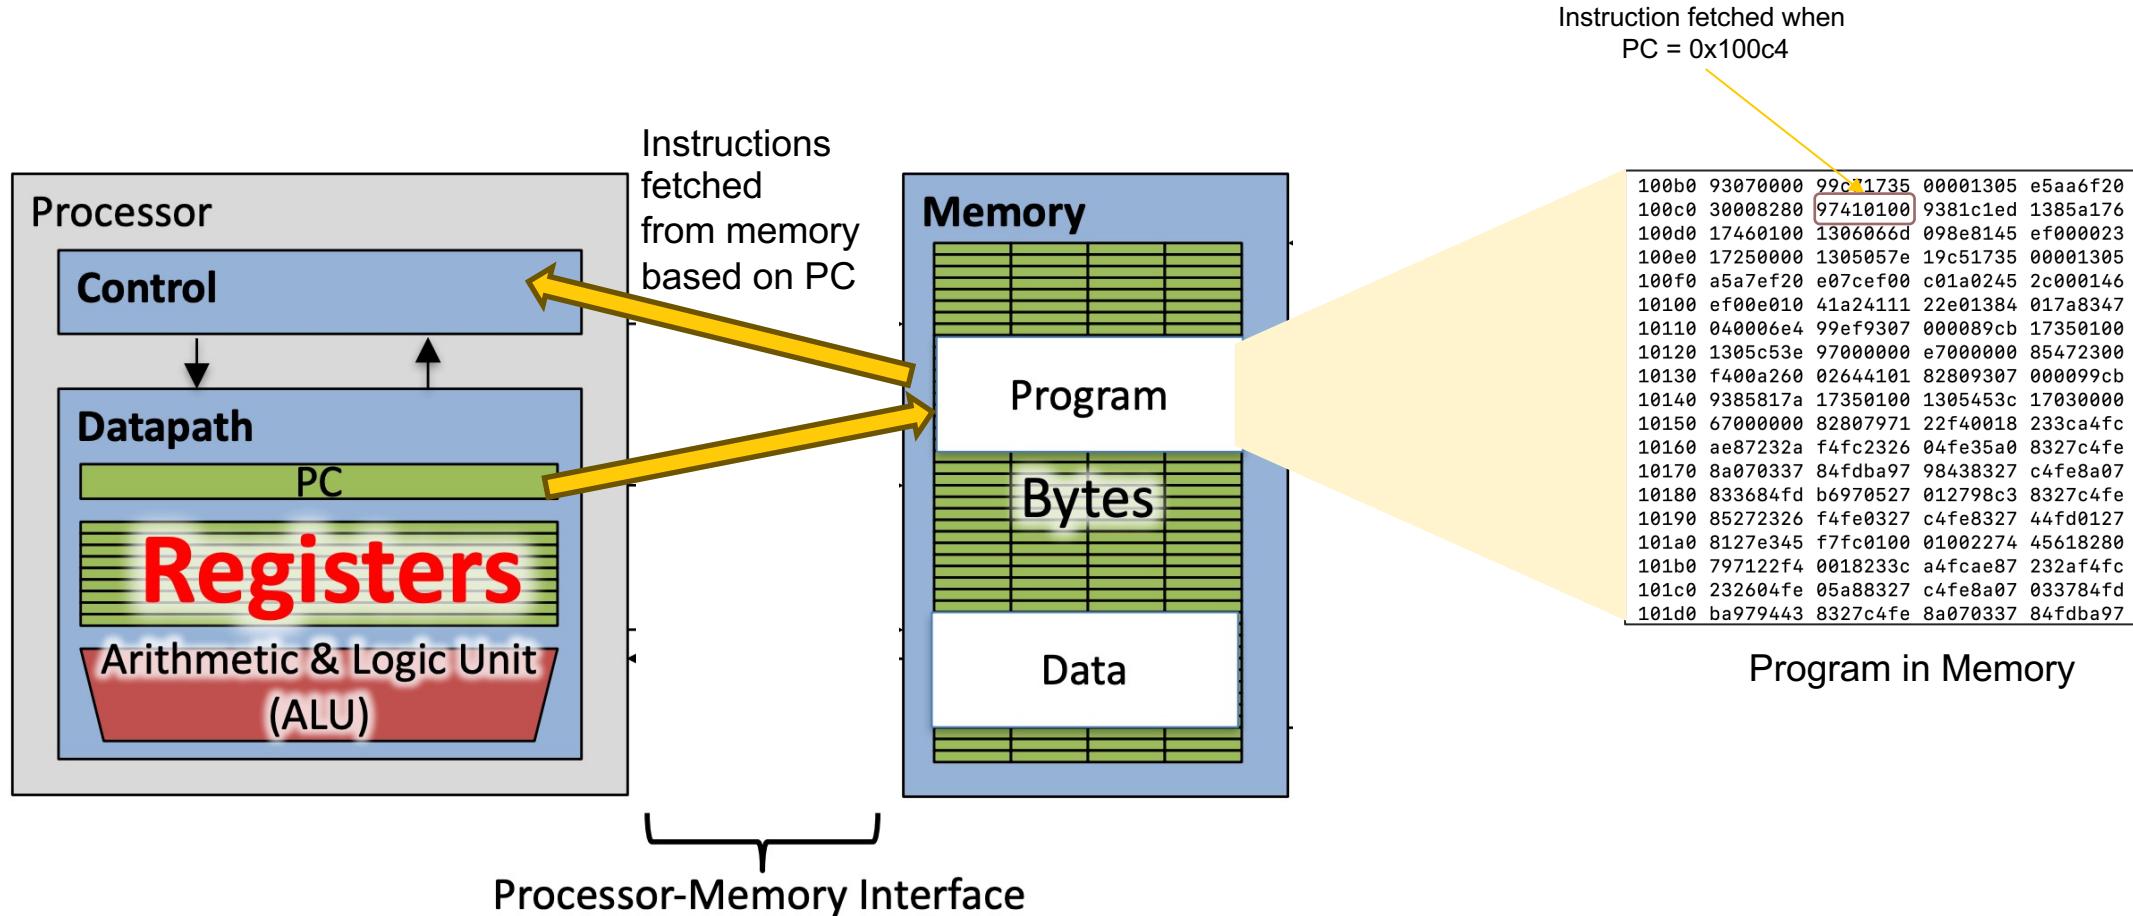
```
x2 holds A17
```

```
lhu x1, 0x0(x2)
```

Memory

A0	A1	A2	A3
A4	A5	A6	A7
A8	A9	A10	A11
A12	A13	A14	A15
A16	A17	A18	A19
A20	A21	A22	A23
A24	A25	A26	A27
A28	A29	A30	A31

The Program Counter (pc)



PC: Points to the current instruction being executed. Typically increments in 4 bytes. However, the sequence breaks due to branches.

Conditional Branch Instructions

offset
with respect
to the PC
(PC relative addressing)

C conditional statement

```
if (a != b) {
    // execute here if TRUE
}
```

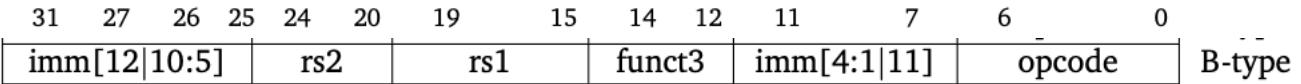
Addr:

```
bne x4, x5, label
// execute instructions
// here if x4 == x5
```

label:

...

Address to branch:
target=PC + label



Inst	Name	FMT	Opcode	funct3	funct7	Description (C)	Note
beq	Branch ==	B	1100011	0x0		if(rs1 == rs2) PC += imm	
bne	Branch !=	B	1100011	0x1		if(rs1 != rs2) PC += imm	
blt	Branch <	B	1100011	0x4		if(rs1 < rs2) PC += imm	
bge	Branch ≥	B	1100011	0x5		if(rs1 >= rs2) PC += imm	
bltu	Branch < (U)	B	1100011	0x6		if(rs1 < rs2) PC += imm	zero-extends
bgeu	Branch ≥ (U)	B	1100011	0x7		if(rs1 >= rs2) PC += imm	zero-extends

Branch range: +/- 4KiB for beq, bne, blt, bge;
+8KiB for bltu and bgeu

Can only jump to even addresses.



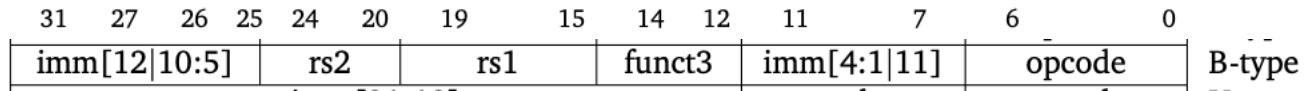
Conditional Branch Instructions

Addr:

```
bne x4, x5, label
// execute instructions
// here if x4 == x5
```

label:

```
. . .
. . .
. . .
```



```
If branch is not taken{
    PC = PC + 4
} else if branch is taken{
    PC = PC + 4 * Immediate
}
```

immediate is number of instructions to jump (remember, specifies words)
either forward (+) or backwards (-)



Pseudo Instructions

- Easy to understand and use. They translate to base instructions
- Example 1: Copy register from rs1 to rd.
Pseudo instruction **mv rd, rs1** translates to **addi rd, rs1, 0**
- Example 2: Load immediate into registers
Pseudo instruction **li rd, CONSTANT**
- Example 3: Load address of symbol
Pseudo instruction **la rd, SYMBOL**
translates to
`lui rd, SYMBOL[31:12]
addi rd, t0, SYMBOL[11:0]`

RISC V Assembly Program Manual: <https://shakti.org.in/docs/risc-v-asm-manual.pdf>

RISC V Assembly Programmer's Manual: <https://github.com/riscv-non-isa/riscv-asm-manual/blob/master/riscv-asm.md>

RISC V Programming Book: <https://riscv-programming.org/book/riscv-book.html>



First RISC V Assembly Program

Strlen C code

```
char str[] = "Hello World!"  
int i;  
for (i=0; str[i] != '\0'; i);  
return i
```

strlen.S

```
.section .text          # represents the section for the code  
.global strlen  
.global _start  
  
.start:                 # execution starts from this label  
strlen:  
    la      a1, str      # load address of str into a1  
    li      t0, 0         # i = 0  
1: # Start of for loop  
    add    t1, t0, a1     # Add the byte offset for str[i]  
    lb     t1, 0(t1)       # Dereference str[i]  
    beqz   t1, 1f         # if str[i] == 0, break for loop  
    addi   t0, t0, 1       # Add 1 to our iterator  
    j      1b             # Jump back to condition (1 backwards)  
1: # End of for loop  
    mv     a0, t0          # Move t0 into a0 to return  
    ret                  # Return back via the return address register  
  
.data                   # represents the section for the data  
str:  .ascii "Hello World!"
```

First RISC V Assembly Program

strlen.S

```
.section .text          # represents the section for the code
.global strlen          # represents the global symbol
.global _start           # represents the entry point

_start:
_strlen:
    la    a1, str      # load address of str into a1
    li    t0, 0          # i = 0

1: # Start of for loop
    add   t1, t0, a1    # Add the byte offset for str[i]
    lb    t1, 0(t1)     # Dereference str[i]
    beqz t1, 1f          # if str[i] == 0, break for loop
    addi  t0, t0, 1      # Add 1 to our iterator
    j     1b             # Jump back to condition (1 backwards)

1: # End of for loop
    mv    a0, t0          # Move t0 into a0 to return
    ret                  # Return back via the return address register

.data
str:  .ascii "Hello World!" # represents the section for the data
```

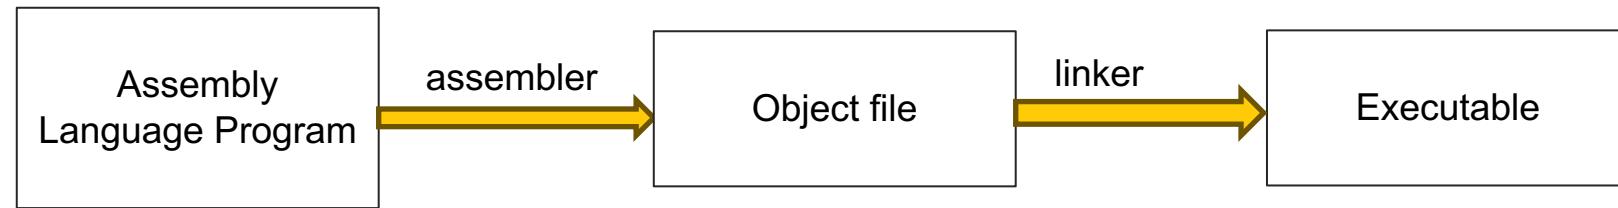
Assembly Directives.

Starts with a '.' used to control the assembler

Labels are markers used to represent program locations

These are pseudo-instructions

Create an Executable



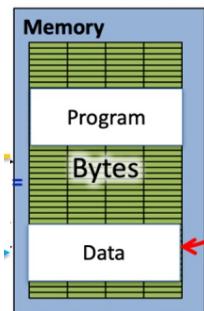
```
riscv64-unknown-elf-gcc -nostdlib -nostartfiles -T ./spike.lds strlen.S -o strlen.elf
```

Invokes the assembler

Linker Descriptor File
 How should the executable be placed in memory?
 The memory map:

Input assembly program

Output executable





The Linker Descriptor File

```
OUTPUT_ARCH("riscv")

ENTRY(_start) #execution should start
              from this label

SECTIONS
{
    . = 0x80000000; #at address

    .text : ALIGN(4K) #place code section
    {
        *(.text)
        *(.text.init)
    }

    .data : ALIGN(4K) #place data section
    {
        *(.data)
    }

}
```



The Linker Descriptor File

```
OUTPUT_ARCH("riscv")

ENTRY(_start) #execution should start
              from this label

SECTIONS
{
    . = 0x80000000; #at address

    .text : ALIGN(4K) #place code section
    {
        *(.text)
        *(.text.init)
    }

    .data : ALIGN(4K) #place data section
    {
        *(.data)
    }
}
```

Checking out the .text part of the executable

```
root@b32eeece94b2:/home/os-iitm/xv6-riscv# riscv64-unknown-elf-objdump -dj .text strlen.elf

strlen.elf:      file format elf64-littleriscv

Disassembly of section .text:

0000000080001000 <_start>:
  80001000: 00001597          auipc   a1,0x1
  80001004: 00058593          mv      a1,a1
  80001008: 4281              li      t0,0
  8000100a: 00b28333          add    t1,t0,a1
  8000100e: 00030303          lb     t1,0(t1)
  80001012: 00030463          beqz   t1,8000101a <_start+0x1a>
  80001016: 0285              addi   t0,t0,1
  80001018: bfcd              j      8000100a <_start+0xa>
  8000101a: 8516              mv     a0,t0
  8000101c: 8082              ret
```





The Linker Descriptor File

```
OUTPUT_ARCH("riscv")

ENTRY(_start) #execution should start
              from this label

SECTIONS
{
    . = 0x80000000; #at address

    .text : ALIGN(4K) #place code section
    {
        *(.text)
        *(.text.init)
    }

    .data : ALIGN(4K) #place data section
    {
        *(.data)
    }
}
```

Checking out the .data part of the executable

root@b32eeece94b2:/home/os-iitm/xv6-riscv# riscv64-unknown-elf-objdump -sj .data strlen.elf

strlen.elf: file format elf64-littleriscv

Contents of section .data:

80002000 48656c6c 6f20576f 726c6421	Hello World!
-------------------------------------	--------------



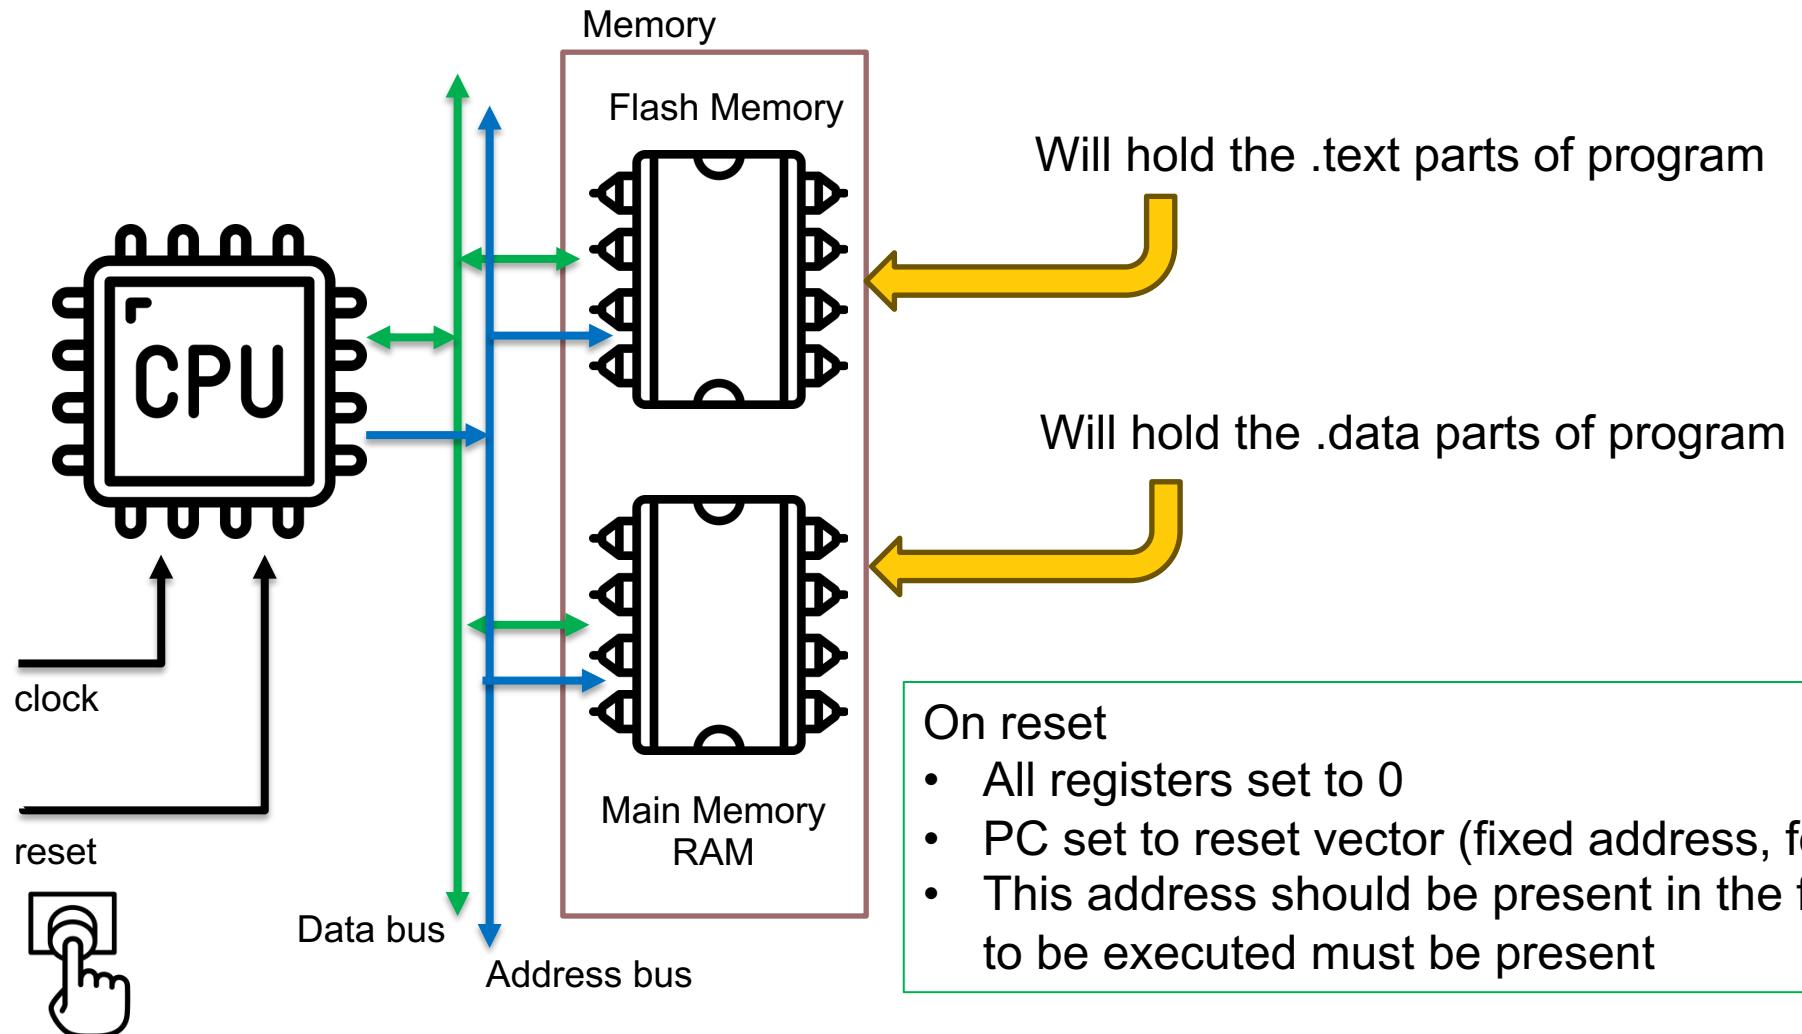


Executing the program

- With an Operating System
 - The OS loads the program into memory based on the memory map
 - It sets the program counter to address of entry (ie. `_start`)
 - Program executes
- Without an Operating System
 - Bare metal

(more on this later)

Executing the program – Baremetal





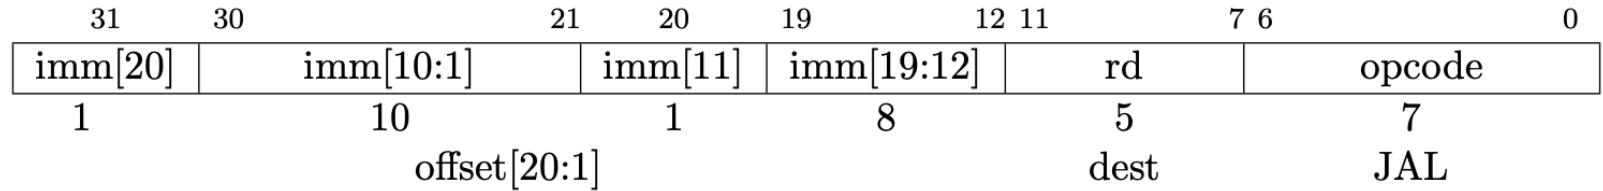
Support for Function Calls

- Invoking functions
- Returning from a function
- Passing parameters
- Managing local variables

Instructions to Invoke a Function Call

jal rd, imm

Jump to $\text{imm} + \text{PC}$, while at the same time store $\text{PC} + 4$ in register rd



Imm is an offset, sign extended and added to PC.

Thus, all jumps are relative to PC.

Jumps can be +/- 1MB from the current PC position

Pseudo instruction `j label` → realized as `jal x0, label`

Indirection Jumps

jalr rd, rs1, imm

Indirect Jump to $rs1+imm$, while at the same time store PC+4 in register rd

31	20 19	15 14	12 11	7 6	0
imm[11:0]	rs1	funct3	rd	opcode	
12 offset[11:0]	5 base	3 0	5 dest	7 JALR	

Note. Jalr is not relative to the PC

Used for function pointers and returns

Pseudo instruction **ret rs1** → realized as **jalr x0, rs1, 0**

If rs1 is not specified, use x1 (ra)

[RISC V Calling Convention](#)

Function Calls (Example)

```

void fact(){
    int f = 1;
    int n = 5;

    while(n > 1){
        f = f * n;
        n = n - 1;
    }
    return;
}

int main(){
    fact();
}
  
```

	0000000000010150 <fact>:		
	10150: 1101	addi	sp, sp, -32
	10152: ec22	sd	s0, 24(sp)
	1018e: 6462	ld	s0, 24(sp)
	10190: 6105	addi	sp, sp, 32
	10192: 8082	ret	
			Sets PC back to 101a0 (based on the ra register)
	1019c: fb5ff0ef	jal	ra, 10150 <fact>
	101a0: 4781	li	a5, 0



More Pseudo Instructions

- Example 4: Unconditional Jumps
Pseudo instruction **j label** translates to **jal x0, label**
Pseudo instruction **jr label** translates to **jalr x0, rs1, 0**
- Example 5: return
Pseudo instruction **ret** translates to **jalr x0, ra, 0**

RISC V Assembly Program Manual: <https://shakti.org.in/docs/risc-v-asm-manual.pdf>

RISC V Assembly Programmer's Manual: <https://github.com/riscv-non-isa/riscv-asm-manual/blob/master/riscv-asm.md>

RISC V Programming Book: <https://riscv-programming.org/book/riscv-book.html>

Multiple function invocations

```
int dec(int n){  
    return n--;  
}
```

```
void fact(){  
    int f = 1;  
    int n = 5;  
  
    while(n > 1){  
        f = f * n;  
        n = dec(n);  
    }  
    return;  
}
```

```
int main(){  
    fact();  
}
```

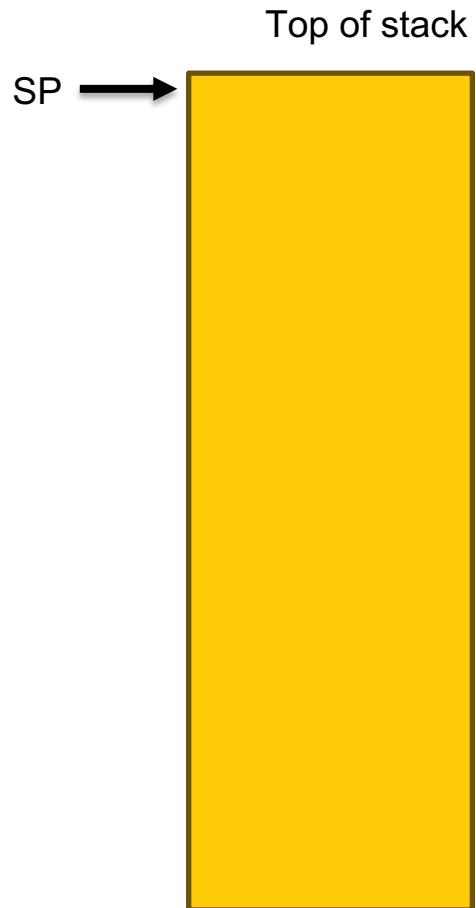
The jal here would over write the ra register

Won't return to main (as ra is overwritten)

The jal here store PC+4 in the ra register

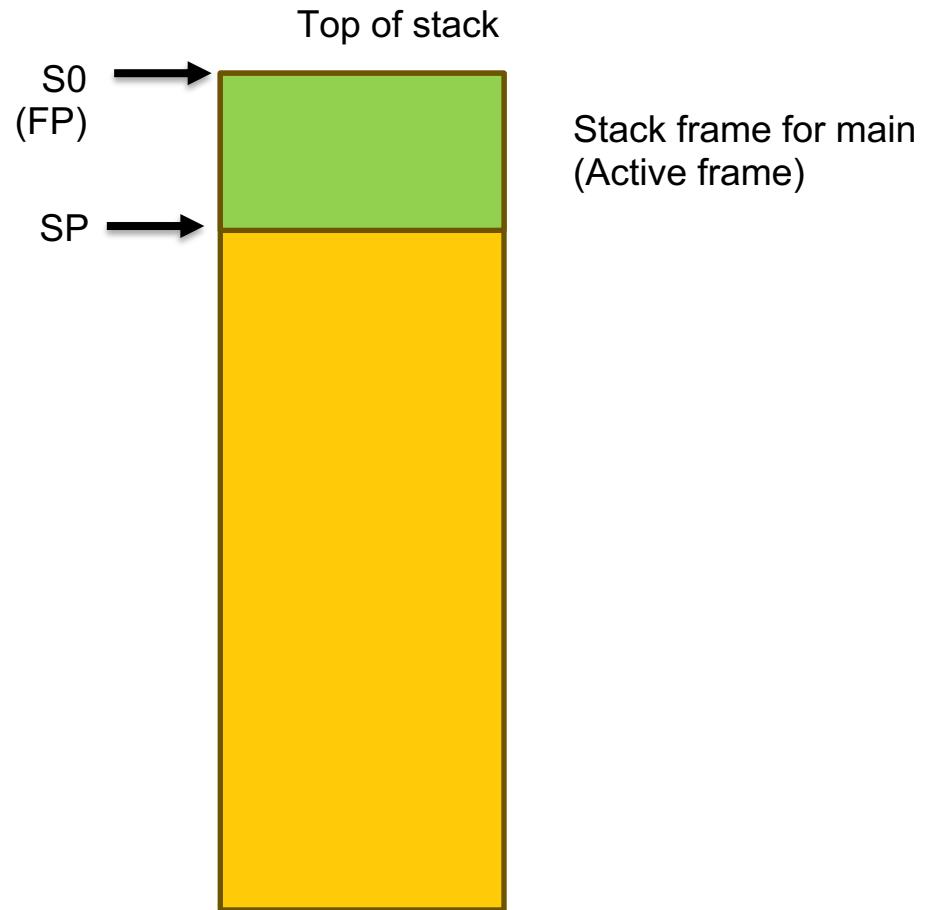
The Stack

- Used to store return addresses
- Also used to store local variables
- Grows from high address to low address
- Register sp (x2) used to point to the lowest address in the stack



Functioning of Stack

Stack grows with each function call.



```
int main(){  
    fact();  
}
```

Corresponding to each function that is executing, there is a stack frame defined between the registers S0 (also called frame pointer (FP)) and stack pointer SP

Functioning of Stack

Stack grows with each function call.

```
void fact(){
    int f = 1;
    int n = 5;

    while(n > 1){
        f = f * n;
        n = dec(n);
    }
    return;
}

int main(){
    fact();
}
```

Prologue: At the start of fact function

addi	sp, sp, -32	Create a stack frame (32 bytes)
sd	ra, 24(sp)	Store the return address in stack
sd	s0, 16(sp)	Store S0 in stack
addi	s0, sp, 32	S0 points to start of stack frame

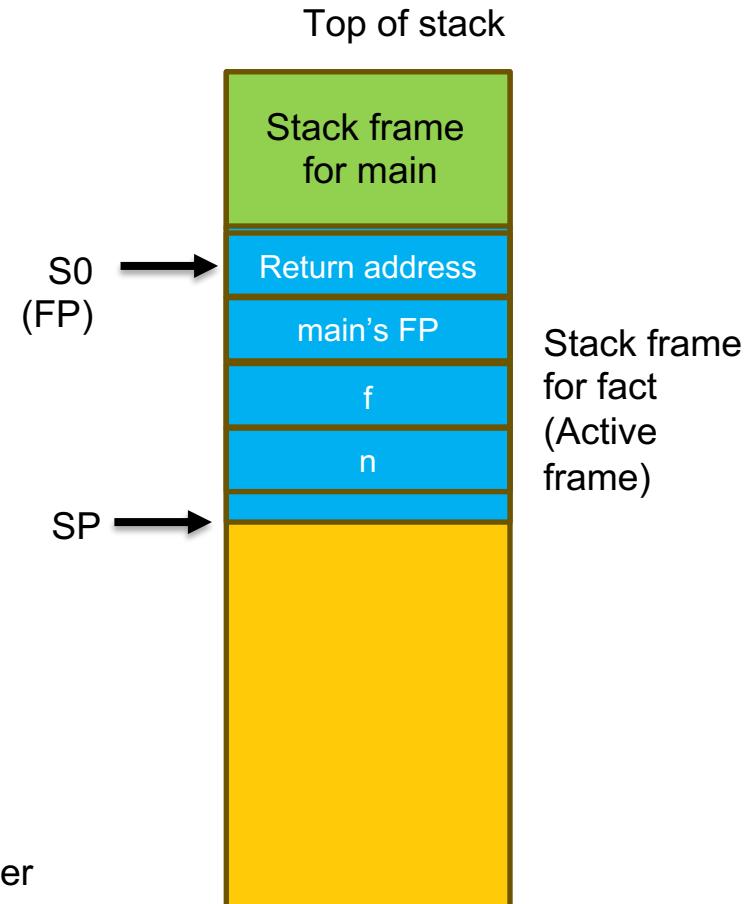
Create the local variables

li	a5, 1
sw	a5, -20(s0) f
li	a5, 5
sw	a5, -24(s0) n

Access local variables

lw	a4, -20(s0)
lw	a5, -24(s0)
mulw	a5, a4, a5

s0 (aka FP): Frame pointer
sp : stack pointer



The stack frame of a function is used to store the previous frame pointer, return address local variables, and any temporary data that needs to be saved.

Functioning of Stack

```

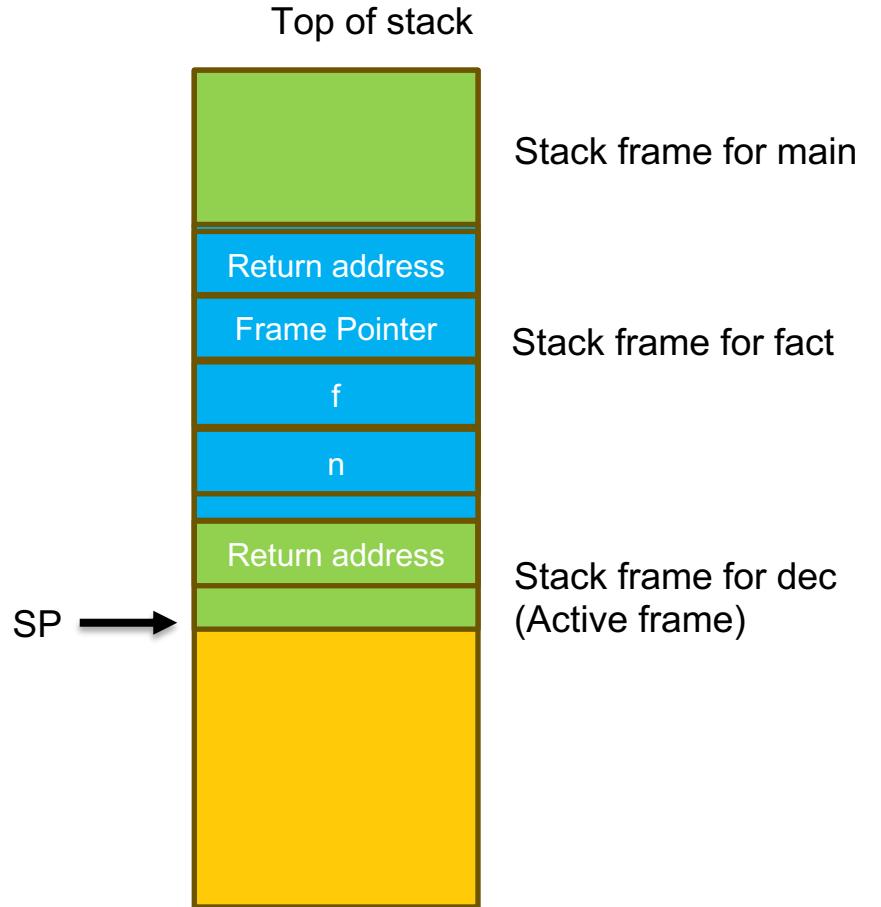
→ int dec(int n){
    return n--;
}

void fact(){
    int f = 1;
    int n = 5;

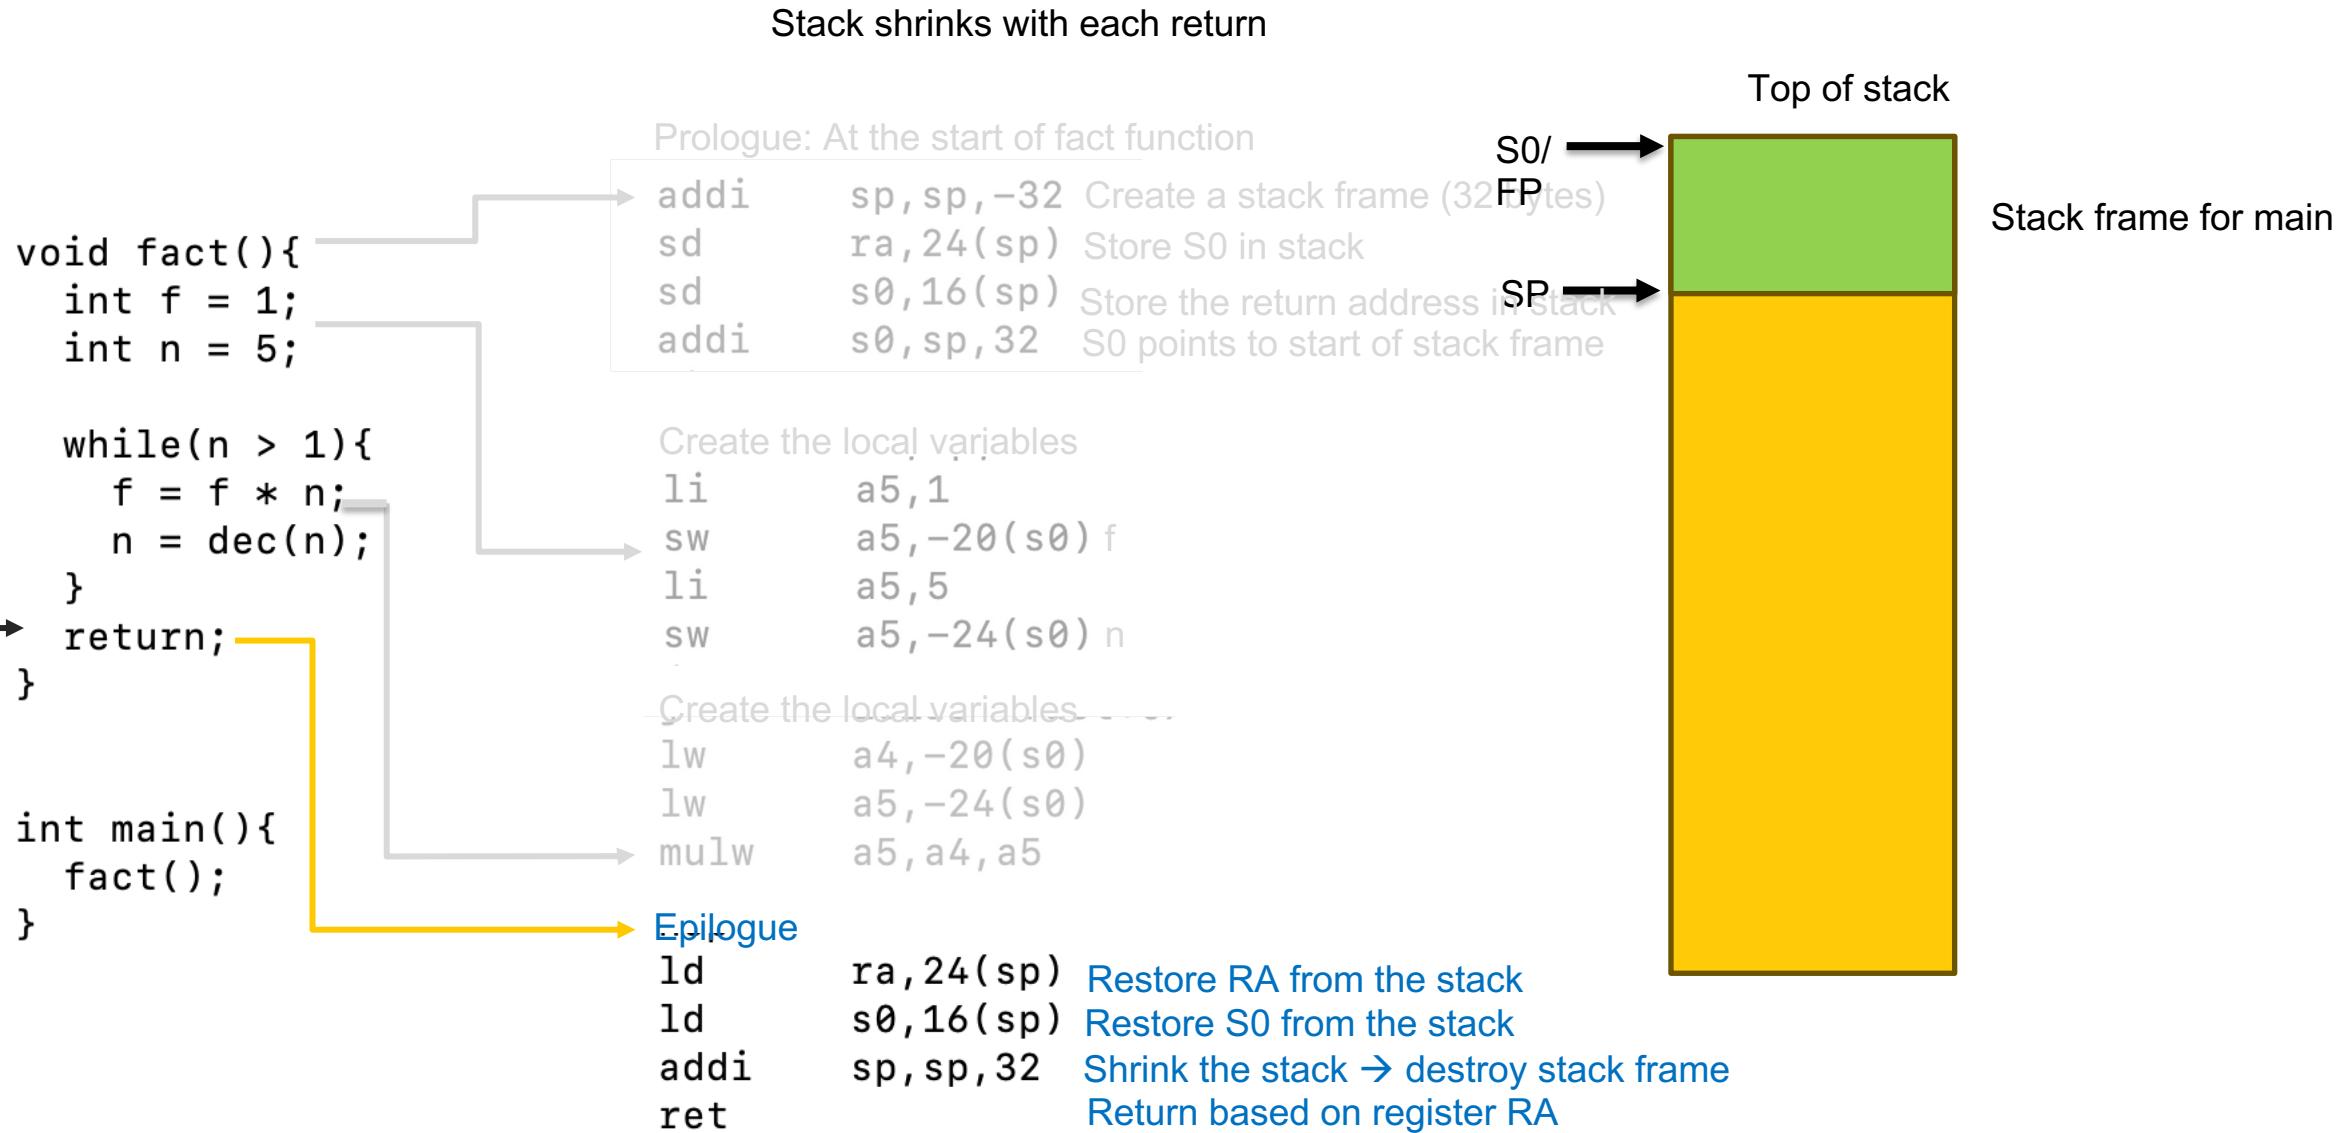
    while(n > 1){
        f = f * n;
        n = dec(n);
    }
    return;
}

int main(){
    fact();
}
  
```

Stack grows with each function call.

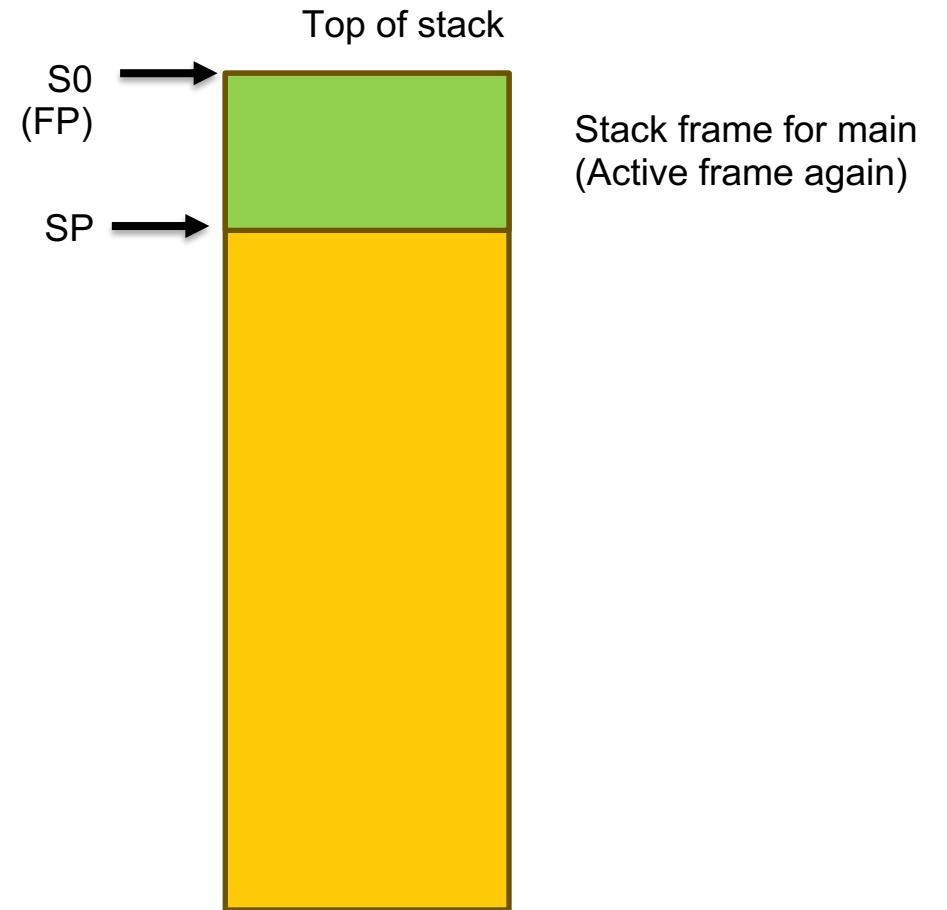


Functioning of Stack



Functioning of Stack

Stack shrinks with each return



```
int main(){  
    fact();  
} → }
```

Passing parameters to functions

- Parameters passed through registers

x10–11	a0–1	Function arguments/return values	Caller
x12–17	a2–7	Function arguments	Caller

```
int sumup(int a, int b, int c, int d, int e,
          int f, int g, int h, int i, int j){
    int res;

    res = a + b + c + d + e + f + g + h + i + j;
    return res;
}
```

- a, b, c, d, e, f, g, h → passed through registers a0 to a7
- i, j → passed through the stack

sumup(1, 2, 3, 4, 5, 6, 7, 8, 9, 10);

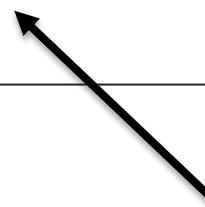


Return data from function

- Return through a0 register

```
int sumup(int a, int b, int c, int d, int e,
          int f, int g, int h, int i, int j){
    int res;

    res = a + b + c + d + e + f + g + h + i + j;
    return res;
}
```



Result returned through a0 register



Example of Parameter Passing and Return

```
int main() {
    int a0 = 5;
    int a1 = 7;

    int result = sum_function(a0, a1);

    return 0;
}

int sum_function(int a, int b) {
    return a + b;
}
```

```
.section .text
.global _start

_start:
    # Example parameters
    li  a0, 5          # First parameter
    li  a1, 7          # Second parameter

    jal ra, sum_function

    ret

.global sum_function
sum_function:
    # Sum the parameters a0 and a1
    add  a0, a0, a1

    ret
```

The two parameters a0 and a1 are passed through resp. registers
The return from sum_function is in the a0 register

Caller and Callee Saved Registers

	Register	ABI Name	Description	Saver
Numbers hardware understands	x0	zero	Hard-wired zero	—
	x1	ra	Return address	Caller
	x2	sp	Stack pointer	Callee
	x3	gp	Global pointer	—
	x4	tp	Thread pointer	—
	x5	t0	Temporary/alternate link register	Caller
	x6–7	t1–2	Temporaries	Caller
	x8	s0/fp	Saved register/frame pointer	Callee
	x9	s1	Saved register	Callee
Human-friendly symbolic names in assembly code	x10–11	a0–1	Function arguments/return values	Caller
	x12–17	a2–7	Function arguments	Caller
	x18–27	s2–11	Saved registers	Callee
	x28–31	t3–6	Temporaries	Caller

Caller Saved (volatile / call-clobbered)

- Holds temporary data used in functions
- Caller function's responsibility to push these data onto the stack before calling any function
- These values will be restored when the function returns

```
someFunction () {
    .
    .
    someOtherFunction ();
    .
    .
}
```

Save all the caller saved registers that you use in this function

Restore caller saved registers here (only if needed)

Caller and Callee Saved Registers

	Register	ABI Name	Description	Saver
Numbers hardware understands	x0	zero	Hard-wired zero	—
	x1	ra	Return address	Caller
	x2	sp	Stack pointer	Callee
	x3	gp	Global pointer	—
	x4	tp	Thread pointer	—
	x5	t0	Temporary/alternate link register	Caller
	x6–7	t1–2	Temporaries	Caller
	x8	s0/fp	Saved register/frame pointer	Callee
	x9	s1	Saved register	Callee
Human-friendly symbolic names in assembly code	x10–11	a0–1	Function arguments/return values	Caller
	x12–17	a2–7	Function arguments	Caller
	x18–27	s2–11	Saved registers	Callee
	x28–31	t3–6	Temporaries	Caller

Callee Saved (non-volatile / call-preserved)

- Holds temporary data used in functions
- Callee function's responsibility to push these data onto the stack at the start of the function
- These values will be restored when the function ends

someOtherFunction () {

• _____

• _____

• _____

• _____

}

Save all the callee saved registers on stack that you use in this function

Restore callee saved registers here



Example for Caller vs Callee Saved Registers

```
.global caller_saved_demo

caller_saved_demo:

    # Function body
    add  t0, a0, a1
    add  t1, t0, a2

    ret
```

Callee saved registers like t0 and t1 do not need to be saved and restored in the function.

```
.global callee_saved_demo

callee_saved_demo:
    # Save callee-saved registers
    add  sp, sp, -4
    sw   s0, 0(sp)
    add  sp, sp, -4
    sw   s1, 0(sp)

    # Function body
    add  s0, a0, a1
    add  s1, s0, a2

    # The result is stored in s1

    # Restore callee-saved registers
    lw   s1, 0(sp)
    add sp, sp, 4
    lw   s0, 0(sp)
    add sp, sp, 4

    ret
```

Caller saved registers like s0 and s1 saved and restored in the function before and after use
Resp.

More examples of Caller saved register

```
caller_saved_demo:

# Function body
add t0, a0, a1

ret
```

The function modifies t0 register, which is a callee saved register

```
_start:
# Initialize arguments
li a0, 5
li a1, 7
li a2, 10

# Call the caller-saved function
jal ra, caller_saved_demo

ret
```

At the caller, no need to save t0, because it is not used in the function

```
_start:
# Use t0 for some computation
li t0, 10
add t0, t0, a0

# Save the value of t0 onto the stack
addi sp, sp, -4
sw t0, 0(sp)

# Call the callee function
jal ra, caller_function_demo

# Restore the value of t0 from the stack
lw t0, 0(sp)
addi sp, sp, 4

# Continue in the caller function if needed
# ...

ret
```

At the caller, save t0 on the stack because it is used in the function. Restore after the caller_function_demo is invoked.